Certifying Graph-Manipulating C Programs via Localizations within Data Structures

ANONYMOUS AUTHOR(S)

We develop powerful and general techniques to mechanically verify realistic programs that manipulate heap-represented graphs and related data structures with intrinsic sharing. We construct a modular and general setup for reasoning about abstract mathematical graphs and use separation logic to define how such abstract graphs are represented concretely in the heap. We upgrade Hobor and Villard's theory of ramification to support existential quantifiers in postconditions and to smoothly handle modified program variables. We demonstrate the generality and power of our techniques by integrating them into the Verified Software Toolchain and certifying the correctness of six graph-manipulating programs written in CompCert C, including a 400-line generational garbage collector for the CertiCoq project. While doing so, we identify two places where the semantics of C is too weak to define generational garbage collectors of the sort used in the OCaml runtime. Our proofs are entirely machine-checked in Coq.

1 INTRODUCTION

Over the last fifteen years, separation logic has facilitated great strides in verifying programs that manipulate tree-shaped data structures. [Appel et al. 2014; Bengtson et al. 2012; Berdine et al. 2005; Chin et al. 2010; Chlipala 2011; Jacobs et al. 2011]. Unfortunately, programs that manipulate graph-shaped data structures (i.e. structures with *intrinsic sharing*) have proved harder to verify. Indeed, such programs were formidable enough that a number of the early landmark results in separation logic devoted substantial effort to verifying single examples such as Schorr-Waite [Yang 2001] or union-find [Krishnaswami 2011] with pen and paper. More recent landmarks have moved to a machine-checked context, but have still been devoted to either single examples or to classes of closely-related examples such as garbage collectors [Ericsson et al. 2017; McCreight et al. 2010]. These kinds of examples tend to require a large number of custom predicates and subtle reasoning, which generally does not carry to the verification of other graph-manipulating programs.

In contrast, we present a general toolkit for verifying graph-manipulating programs in a machine-checked context. Our techniques are *general* in that they handle a diverse range of graph-manipulating programs, and *modular* in that they allow code reuse (*e.g.* facts about reachability) and encourage separation of concerns (*e.g.* between abstract mathematical graphs and their concrete representation in the heap). Our techniques are *powerful* enough to reason about real C code as compiled by CompCert [Leroy 2006], and also *lightweight* enough to integrate into the Verified Software Toolchain (VST) [Appel et al. 2014] without requiring major reengineering. Both CompCert and VST are distributed as optional packages in the CoqIDE installer and through opam, so they have a sizable user base that can take advantage of our techniques. Finally, our techniques *scale* well beyond short toy programs: we certify the correctness of a generational garbage collector for the CertiCoq project [Anand et al. 2017] (≈ 400 rather devilish lines of C).

We proceed in three steps. First, we develop a "mathematical graph library" that is general enough to reason about a wide variety of algorithms and expressive enough to describe the behavior of these algorithms in real machines. We modularize this library carefully so that common ideas—*e.g.* subgraphs, reachability, and isomorphism—can be efficiently reused in different algorithms. Second, use separation logic to express how these abstract graphs are actualized in the heap as concrete graphs in a way that facilitates the reuse of key definitions and theorems across algorithms. Finally,

we develop a notion of *localization blocks* that allows us to carry out our Hoare proofs in a modular fashion, even in the presence of the implicit sharing intrinsic to graphs, by using our Localize rule:

$$
\frac{G_1 \vdash L_1 * R \qquad \{L_1\}\ c\ \{\exists x.\ L_2\} \qquad R \vdash \forall x.\ (L_2 \mathbin{-\!\!*} G_2)}{\{G_1\}\ c\ \{\exists x.\ G_2\}} \quad \mathrm{FreeVar}(R) \cap \mathrm{ModVar}(c) = \emptyset \tag{1}
$$

Localize connects the "local" effect of a command c, *i.e.* transforming L_1 to L_2, with its "global" effect, *i.e.* from G_1 to G_2. The key is carefully choosing a *ramification frame R* that satisfies a pair of delicately-stated entailments[1] and the side condition on modified local program variables. Localize is a more general version of the well-known Frame rule, which does the same task in the simpler case when $G_i = L_i * F$ for some frame F that is untouched by c. Said differently, Frame works well for tree-manipulating programs, while Localize handles the more subtle graph-manipulating programs. Localize upgrades the Ramify rule [Hobor and Villard 2013] in two key ways: support for existential quantifiers in postconditions and smoother treatment of modified program variables.

Our contributions are organized as follows:

§2 We use the classic "union-find" disjoint set algorithm [Cormen et al. 2009] to show how our three key ingredients—mathematical graphs, spatial graphs, and localization blocks—come together to verify graph-manipulating algorithms. To the best of our knowledge this is the first machine-checked verification of this algorithm that starts with real C code. We introduce *localization blocks* as a notation for Localize in decorated programs.

§3 We show that Localize and Frame are equivalent, and a delicate technique to properly handle modified local variables. We show a mark-graph program that explores a graph in a fold/unfold style and discuss the utility of linked existentials in postconditions. We also briefly discuss some additional examples to give a sense of the breadth of algorithms we can verify: marking a DAG, an array-based version of union-find, and pruning a graph into a spanning tree.

§4 We develop a general framework of mathematical graphs powerful enough to support realistic verification in a mechanized context. We give a sampling of key definitions and show how our framework is modularized to facilitate code reuse.

§5 We suggest that the Knaster-Tarski fixpoint [Tarski 1955] cannot define a usable separation logic graph predicate. We propose a better definition for general spatial graphs that still enjoys a "recursive" fold/unfold. We prove general theorems about spatial graphs in a way that can be utilized in multiple flavors of separation logic.

§6 We discuss the certification of the CertiCoq garbage collector (GC). While certifying this GC we identify two places where the semantics of C is too weak to define an OCaml-style GC. We also found and fixed a rather subtle overflow error in the original C code for the GC, justifying the effort of developing the machine-checked proof.

§7 We discuss how our techniques are integrated into the "Floyd" module of VST, a separation-logic based engine to help users verify CompCert C programs, via two new Floyd tactics `localize` and `unlocalize`. We also document statistics related to our overall development.

§8 We discuss related work.

§9 We discuss directions for future work and conclude.

All of our results are machine checked in Coq and available at [johndoe20190406 2019].

---

[1]Readers less familiar with the separating implication $P \mathbin{-\!\!*} Q$, also known as *magic wand*, can refer to its semantics in Figure 7 (page 13), which also models the other separation logic operators we use in this paper. The key proof rule for magic wand is its adjointness with the separating conjunction: $(P * Q \vdash R) \Leftrightarrow (P \vdash Q \mathbin{-\!\!*} R)$.

```
1  struct Node { unsigned int rank;
2                struct Node * parent; }
3  // {uf_graph(γ) ∧ x ∈ V(γ)}
4  struct Node* find(struct Node* x) {
5    struct Node *p;
```

6 $// \searrow \begin{cases} \text{uf\_graph}(\gamma) \land x \in V(\gamma) \land \\ \exists r, pa. \ \gamma(x) = (r, pa) \land pa \in V(\gamma) \end{cases}$

7 $// \searrow \begin{cases} x \mapsto r, pa \land x \in V(\gamma) \land \\ \gamma(x) = (r, pa) \land pa \in V(\gamma) \end{cases}$

```
8  ↯(2) p = x -> parent;
```

9 $// \nearrow \begin{cases} x \mapsto r, pa \land \text{p} = pa \land x \in V(\gamma) \land \\ \gamma(x) = (r, pa) \land pa \in V(\gamma) \end{cases}$

10 $// \begin{cases} \text{uf\_graph}(\gamma) \land \text{p} = pa \land x \in V(\gamma) \land \\ \gamma(x) = (r, pa) \land pa \in V(\gamma) \end{cases}$

```
11   if (p != x) {
```

12 $// \begin{cases} \text{uf\_graph}(\gamma) \land \text{p} = pa \land pa \neq x \land \\ x \in V(\gamma) \land \gamma(x) = (r, pa) \land pa \in V(\gamma) \end{cases}$

```
13     p = find(p);
```

14 $// \begin{cases} \exists \gamma', rt. \ \text{uf\_graph}(\gamma') \land \text{p} = rt \land pa \neq x \land x \in V(\gamma) \land \\ findS(\gamma, pa, \gamma') \land uf\_root(\gamma', pa, rt) \land \gamma(x) = (r, pa) \end{cases}$

15 $// \searrow \begin{cases} x \mapsto r, pa \land \text{p} = rt \land pa \neq x \land findS(\gamma, pa, \gamma') \land \\ uf\_root(\gamma', pa, rt) \land x \in V(\gamma) \land \gamma(x) = (r, pa) \end{cases}$

```
16   ↯(3) x -> parent = p;
```

17 $// \nearrow \begin{cases} x \mapsto r, rt \land \text{p} = rt \land pa \neq x \land findS(\gamma, pa, \gamma') \land \\ uf\_root(\gamma', pa, rt) \land x \in V(\gamma) \land \gamma(x) = (r, pa) \end{cases}$

18 $// \begin{cases} \exists \gamma''. \ \text{uf\_graph}(\gamma'') \land findS(\gamma, pa, \gamma'') \land \\ uf\_root(\gamma'', x, rt) \land \text{p} = rt \end{cases}$

```
19   } return p;
```

20 $\}$ $// \begin{cases} \exists \gamma'', rt. \ \text{uf\_graph}(\gamma'') \land findS(\gamma, x, \gamma'') \land \\ uf\_root(\gamma'', x, rt) \land \text{ret} = rt \end{cases}$

$$\text{uf\_graph}(x, \gamma) \stackrel{\Delta}{=} \mathop{\bigstar}_{v \in V(\gamma)} v \mapsto \gamma(v)$$

$$\text{uf\_root}(\gamma, x, rt) \stackrel{\Delta}{=} x \overset{\gamma}{\rightsquigarrow}{}^\star rt \land \\ \forall rt'. \ rt \overset{\gamma}{\rightsquigarrow}{}^\star rt' \Rightarrow rt = rt'$$

$$findS(\gamma, x, \gamma') \stackrel{\Delta}{=} \big(\forall v. \ v \in V(\gamma) \Leftrightarrow v \in V(\gamma')\big) \land$$
$$\big(\forall v. \ v \in V(\gamma) \Rightarrow \gamma(v).rank = \gamma'(v).rank\big) \land$$
$$\big(\forall r, r'. \ uf\_root(\gamma, v, r) \Rightarrow uf\_root(\gamma', v, r') \Rightarrow r = r'\big) \land$$
$$\big(\gamma \smallsetminus \{v \in \gamma \mid x \overset{\gamma}{\rightsquigarrow}{}^\star v\} \cong \gamma' \smallsetminus \{v \in \gamma \mid x \overset{\gamma}{\rightsquigarrow}{}^\star v\}\big)$$

Fig. 1. Clight code and proof sketch for find

2 LOCALIZATIONS YIELD A TIDY UNION-FIND

As an initial demonstration of our techniques, we show the decorated code of the `find` function from the classic disjoint-set data structure in Figure 1. The `find` function returns the root (ultimate parent) of a `Node` x. A node is a root when its parent pointer points to itself (line 11); other than such self-loops at roots, the structure is acyclic. For good amortised performance, `find` also performs path compression (line 16). At first, `find` appears rather trivial since it only has about 5 lines of code and a `Node` has only a single outgoing pointer. In actual fact, the rather subtle nature of path compression and the implicit sharing inherent in parent-pointers make the disjoint-set data structure very difficult to reason about. Indeed, the first pen-and-paper verification in separation logic required 20 pages [Krishnaswami 2011].

We use the following conventions in our invariants. Pure predicates are written in *italic*. We write γ to mean a "mathematical" (or "pure") graph: roughly, a set of labeled vertices $V(\gamma)$ and edges $E(\gamma)$. When $v \in V(\gamma)$, we write $\gamma(v) = (r, p)$ to state that vertex v has label r and parent vertex p (r stores the "rank" of a node; it is ignored in `find`). We detail mathematical graphs in §4.

Spatial predicates are written in sans-serif. Each node $v \in V(\gamma)$ is represented in the heap by $v \mapsto \gamma(v)$, where we use the usual pen-and-paper trick of writing *e.g.* $v \mapsto r, p$ to mean $(v \mapsto r) * ((v + sizeof(\text{unsigned int})) \mapsto p)$ in the character-addressed C memory model. The whole graph (disjoint-set forest) is represented by uf_graph(γ), essentially the iterated separating conjunction of the representations of each vertex $v \in V(\gamma)$. We detail spatial graphs in §5.

The invariants at each program point are natural despite only minor tidying from our machine-checked proof. We also enjoy good separation between the spatial predicates and pure predicates. All of this is despite verifying real C code, which entails quite a number of grungy details. As one example, we will examine some grunginess that occurs in the verification of line 11 shortly.

The precondition on line 3 says that we have a disjoint-set forest representing the abstract graph γ, and that x is a valid vertex in γ. The postcondition is on line 20: the heap contains a new

union-find graph γ'', and `find` returns the node rt. We specify that rt is the root (ultimate parent) of x with the pure relation uf_root. The relation $findS$, which conservatively approximates the action of path compression, relates the final graph γ'' to the original graph γ. The formal definitions for the concepts used in uf_root and $findS$ will be given in §4, but briefly: $x \overset{\gamma}{\rightsquigarrow}^\star y$ expresses that y is reachable from x in γ, $\gamma \setminus S$ expresses the result of removing the vertices in set S from graph γ, and $\gamma_1 \cong \gamma_2$ expresses that the two graphs are structurally equivalent.

Most of the verification is straightforward. To aid human readability, we use the color red to indicate the changes in the invariants line-by-line. Each individual line of code (8, 11, 13, 16, and 19) is bracketed with invariants leading to relatively easy proofs of the command (ignoring the symbols $\searrow$, $\frac{\ell}{\ell}(i)$, and $\swarrow$ until §2.1). In addition to improving human comprehensibility, this also aids mechanical comprehensibility—that is, straightforward invariants help the underlying verification engine (VST, in our case) handle many grungy details for us either automatically or with a little human guidance via suitable lemmas.

The pointer comparison in line 11 is an example of where such lemmas are necessary. Formally, pointer (in-)equality comparison in C is only defined under somewhat delicate circumstances[2]. VST could prove the definedness of the pointer comparison automatically if we knew $(x \mapsto \_) * (p \mapsto \_) * \top$, but unfortunately this does not follow from line 9 since, when x is a root, self-loop gives us $x = p$ and $x \mapsto \_ * x \mapsto \_ \vdash \bot$. Accordingly, we must prove a simple lemma that states that when $uf\_graph(\gamma) \wedge x \in V(\gamma) \wedge p \in V(\gamma)$, the pointer comparison is defined in C.

2.1 Localization Blocks

It is time to explain the non-obvious jumps in reasoning bracketed by the symbols $\searrow$, $\frac{\ell}{\ell}(i)$, and $\swarrow$ (lines 6–10 and 14–18). We call such bracketed sets of lines "localization blocks". As explained above, the verification of the command itself is entirely straightforward given its immediate neighbors (lines 7–9 and 15–17). What is not so straightforward is how *e.g.* line 6 leads to line 7 or how line 9 leads to line 10. Intuitively, a localization block allows us to zoom in from a larger "global" context to a smaller "local" one, and, after verifying some commands locally and arriving at a local postcondition, to zoom back out to the global context. The $\searrow$ and $\swarrow$ symbols formally indicate an application of the LOCALIZE rule (equation 1 from page 2).

Recall that LOCALIZE connects some "global" pre- and postconditions G_1 and G_2 with some "local" pre- and postconditions L_1 and L_2 using a ramification frame R. The lines adjacent to the $\searrow$ and $\swarrow$ symbols specify G_1 (*e.g.* line 6), L_1 (line 7), L_2 (line 9), and G_2 (line 10). Note that the LOCALIZE rule expects quantifiers in the postconditions L_2 and G_2, but lines 9–10 do not have any. We can overcome this mismatch since $\forall P. (P \dashv\vdash \exists x : unit.P)$ for any x not free in P.

What is not specified explicitly is the ramification frame R, which must satisfy the entailments $G_1 \vdash L_1 * R$ and (again eliding the quantifier) $R \vdash L_2 \mathbin{-\!\!*} G_2$. Here things are a little delicate. To give intuition, it is **almost** enough to choose $R \overset{\Delta}{=} L_2 \mathbin{-\!\!*} G_2$, which makes the second entailment trivial, and reduces the problem to three remaining checks.

The first check is the *ramification entailment*: $G_1 \vdash L_1 * (L_2 \mathbin{-\!\!*} G_2)$. Informally, this asks whether replacing L_1 with L_2 inside G_1 yields G_2. This ramification entailment is a nontrivial proof obligation both because we must prove L_1 is located inside G_1, and because we must prove that "replacing" L_1 with L_2 yields G_2. When we want to refer to the ramification entailment of a localization block in subsequent text we use the symbol $\frac{\ell}{\ell}(i)$ to connect it to relevant equation numbers. Accordingly,

---

[2]To summarize: it is a mess. Specifically, whenever (1) x and p are both null; or when (2) one of them is null and the other has offset between 0 and the size of the memory block into which it is pointing; or when (3) if x and p are from the same memory block, then both of their offsets are between 0 and the size of that block; or when (4) x and p are not in the same memory block and both have offsets between 0 and the size of their respective memory blocks *minus one*.

$\oint(2)$ refers to the ramification entailment associated with lines 6–10:

$$\text{graph}(\gamma) \wedge x \in V(\gamma) \quad \vdash \quad x \mapsto \gamma(x) * \left(x \mapsto \gamma(x) \mathrel{-\!\!*} \text{graph}(\gamma)\right) \tag{2}$$

Here we have isolated the key **spatial** parts of the invariants on lines 6–10. Notice that this lemma is stated for any whole-graph predicate $\text{graph}(\gamma)$, and not merely for the special class of "union-find graphs" $\text{uf\_graph}(\gamma)$ (that *e.g.* have only one outgoing edge per node). That is useful because we use the same lemma to prove similar goals in all of our examples. Indeed, this "unchanged vertex" ramification entailment is used whenever we need to read from a vertex in a graph. In §5 we describe other generic and reusable lemmas that prove other ramification entailments.

The second check is the Hoare proof of the local change from L_1 to L_2. Since lines 7 and 9 are straightforward—indeed, the point of Localize is to make them so—verifying line 8 is easy.

The third check is the side condition on the modified variables. Here we have an irritating problem: the free variables of $R \stackrel{\Delta}{=} L_2 \mathrel{-\!\!*} G_2$ are **not** disjoint from the local variables modified by c. Inspection of lines 8–10 shows that the program variable p **is** modified by c, and **is** free in both L_2 and G_2. This issue is fundamental: the whole point of verifying a read is to know something about the value that has been read. Accordingly, our proof fails when we choose $R \stackrel{\Delta}{=} (L_2 \mathrel{-\!\!*} G_2)$. We will address this problem head-on in §3.2, but for now let us content ourselves with knowing that other than this problem, Localize lets us verify lines 6–10.

The second localization block (lines 14–18) is both easier and harder than the first. It is easier because line 16 does not modify any local program variables, so the side condition is trivially satisfied. Moreover, although line 18 does contain an existential, line 17 does not, and so there is no need to "link" the two associated witnesses. We will discuss this issue in more detail in §3.3, but for now it is enough to choose $R \stackrel{\Delta}{=} L_2 \mathrel{-\!\!*} G_2$ exactly written in lines 17 (for L_2) and 18 (for G_2).

On the other hand, the second localization block is harder than the first because there is more going on spatially. $\oint(3)$ expresses an update to a single node of our graph:

$$\frac{x \in V(\gamma') \qquad \gamma'' = [x \to (r, rt)]\gamma'}{\text{graph}(\gamma') \ \vdash \ x \mapsto \gamma'(x) * \left(x \mapsto \gamma''(x) \mathrel{-\!\!*} \text{graph}(\gamma'')\right)} \tag{3}$$

Here we abuse notation a little bit. The conclusion of the "rule" (actually, lemma) is exactly right and appropriately generic, so spatial ramification lemmas of the kind given in §5 can handle the dirty spatial work for us. However, the second premise uses a notation for "mathematical graph node update" that is customized for union-find graphs, since most graphs have more than a rank and single outgoing edge. More seriously, updating a mathematical graph cannot be done willy-nilly; it is only defined when the properties that restrict the mathematical structure of γ are preserved. For example, in the case of union-find graphs, the graph must be acyclic (other than at roots). In Coq, these properties are carried around via dependent types as will explained in §4.1.

In our proof of find, the bulk of the example-specific effort (as opposed to generic lemmas we reuse in other examples) is showing that this mathematical update can be done properly, *i.e.* from

$$\mathrm{x} \in V(\gamma) \wedge \gamma(\mathrm{x}) = (r, pa) \quad \text{and} \quad \mathrm{x} \neq pa \wedge \mathit{findS}(\gamma, pa, \gamma') \wedge \mathit{uf\_root}(\gamma', pa, rt)$$

we can prove

$$\exists \gamma''.\ \gamma'' = [x \to (r, rt)]\gamma' \wedge \mathit{uf\_root}(\gamma'', \mathrm{x}, rt) \wedge \mathit{findS}(\gamma, \mathrm{x}, \gamma'')$$

This lemma captures the essence of both finding the root and doing path compression: after compressing your parent and finding its root, you can path compress yourself by rerouting your own parent pointer to your (soon-to-be former) parent's root. The existential in the goal is nontrivial exactly because the update $[x \to (r, rt)]\gamma'$ is not always kosher. This lemma requires some effort to prove, but is completely isolated from the grungy details of C.

$$\cfrac{G_1 \vdash L_1 * R \qquad \cfrac{\{L_1\} \; c \; \{L_2\}}{\{L_1 * R\} \; c \; \{(\exists x. \; L_2) * R\}} \; \text{\small Frame} \qquad \cfrac{\cfrac{\vdots}{(\exists x. \; L_2) * (\forall x. \; (L_2 \mathbin{-\!\!*} G_2)) \vdash \exists x. \; G_2} \; \text{\small Tauto}}{(\exists x. \; L_2) * R \vdash \exists x. \; G_2} \; \text{\small Cut}}{\{G_1\} \; c \; \{\exists x. \; G_2\}} \; \text{\small Cons}$$

$$\cfrac{\cfrac{\{P\} \; c \; \{Q\} \qquad Q \vdash \exists x_f. \; Q}{\{P\} \; c \; \{\exists x_f. \; Q\}} \; \text{\small Cons} \qquad F \vdash \forall x_f. \; (Q \mathbin{-\!\!*} (Q * F))}{\cfrac{\{P * F\} \; c \; \{\exists x_f. \; (Q * F)\}}{\{P * F\} \; c \; \{Q * F\}}} \; \text{\small Localize} \qquad \cfrac{\exists x_f. \; (Q * F) \vdash Q * F}{} \; \text{\small Cons}$$

Fig. 2. Proving Localize from Frame, and conversely Frame from Localize

With the second localization block complete, the remainder of the verification is straightforward.

3 LINKING EXISTENTIALS IN LOCALIZATIONS

In this section we (§3.1) first ensure our feet are solidly planted by proving why Localize is sound, and indeed equivalent to Frame. Second (§3.2), we address the bug from §2.1 and show that Localize is indeed strong enough to robustly handle modified local program variables. Third (§3.3), we showcase two additional features of our framework, linked existentials and a fold/unfold style for spatial graphs, by covering the verification of a program that marks a cyclic graph. Finally (§3.4), we briefly discuss some additional examples we have handled. Our flagship example, the garbage collector for the CertiCoq project, will be covered in §6.

3.1 Soundness of Localize

In Figure 2 we put the proof sketches that show that Localize and Frame are equivalent. They require a little care with quantifiers, but are in essence straightforward. In the latter proof set $R \stackrel{\Delta}{=} F$, choose x_f fresh, and range the quantifiers over the unit type. Notice that in both directions the restriction on modified program variables is satisfied: in the first proof, Localize's side condition that $\text{FreeVar}(R) \cap \text{ModVar}(c) = \emptyset$ is exactly what Frame needs; in the second, Frame's side condition that $\text{FreeVar}(F) \cap \text{ModVar}(c) = \emptyset$ is exactly what Localize needs (since $R \stackrel{\Delta}{=} F$). The equivalence between Frame and Localize means that our techniques will be sound in any separation logic.

Notes on notation. Although we do not do so in Figure 1, localization blocks can safely nest. When the ramification entailment is not noteworthy we can omit the $\frac{\Delta}{7}(i)$ reference in pen-and-paper proofs. When we wish to save vertical space we can write $\{G_1\} \searrow \{L_1\}$ and $\{G_2\} \nearrow \{L_2\}$. We also note that since Localize can derive Frame, our notation for localization blocks clarifies pen-and-paper uses of Frame, especially in multi-line contexts with nontrivial F.

3.2 Smoothly handling modified program variables

Consider using Localize to verify the program

```
// {x = 5 ∧ A} ↘ {x = 5 ∧ B}
  ...; x = x + 1; ...;
// {x = 6 ∧ D} ↗ {x = 6 ∧ C}
```

Suppose that other (elided) lines of the program make localization desirable, even though it is overkill for a single assignment. The key issue is that if one sets $R \stackrel{\Delta}{=} L_2 \mathbin{-\!*} G_2$, as we tried to do in §2.1, then the program variable x appears in all four positions in the ramification entailment

$$\overbrace{(x=5 \wedge A)}^{G_1} \vdash \overbrace{(x=5 \wedge B)}^{L_1} * \big(\overbrace{(x=6 \wedge C)}^{L_2} \mathbin{-\!*} \overbrace{(x=6 \wedge D)}^{G_2} \big)$$

For the sake of simplicity, assume that in the above snippet only x is modified and that x does not appear free in A, B, C or D. Let us further assume that, modulo the local variable issue we are trying to solve, the entailment holds. In other words, let us assume that $A \vdash B * (C \mathbin{-\!*} D)$.

Turning to the local variable issue itself, in §2.1 we observed that $L_2 \mathbin{-\!*} G_2$ does **not** ignore the modified program variable x, preventing us from meeting LOCALIZE's side condition[3]. Intuitively, the side condition on LOCALIZE seems to be a bit too strong since it prevents us from mentioning variables in the postconditions that have been modified by code c. As in other cases when life gets tough, what we need is an elegant little dance, and as with most dances one should lead by example.

First, define $\hat{L}_2(x_f) \stackrel{\Delta}{=} (x_f = 6 \wedge C)$ and $\hat{G}_2(x_f) \stackrel{\Delta}{=} (x_f = 6 \wedge D)$, *i.e.* replace the troublesome program variable x in L_2 and G_2 with a harmless fresh metavariable x_f. Next, notice that with a carefully chosen existential quantifier, we can express the original L_2 with the new $\hat{L}_2$ while keeping the troublesome program variable x isolated and shift the above decorated program into the form

```
1    ...; x = x + 1; ...;
2      // {x = 6 ∧ C}
3  ✓  // {∃x_f. (x_f = x) ∧ (x_f = 6 ∧ C)}
4      // {∃x_f. (x_f = x) ∧ (x_f = 6 ∧ D)}}
5      // {x = 6 ∧ D}}
```

Notice that lines 3 and 4 are exactly in the form $\exists x_f.\hat{\hat{L}}_2(x_f)$ and $\exists x_f.\hat{\hat{G}}_2(x_f)$, *i.e.* exactly in the format permitted by LOCALIZE, where $\hat{\hat{L}}_2(x_f) \stackrel{\Delta}{=} (x_f = x) \wedge \hat{L}_2(x_f)$, *i.e.* $(x_f = x) \wedge (x_f = 6 \wedge C)$ and $\hat{\hat{G}}_2(x_f)$ is similar. Now apply LOCALIZE with $R \stackrel{\Delta}{=} \forall x_f.\ \hat{L}_2(x_f) \mathbin{-\!*} \hat{G}_2(x_f)$, *i.e.* $\forall x_f.\ (x_f = 6 \wedge C) \mathbin{-\!*} (x_f = 6 \wedge D)$. By construction, R is free from all program variables modified by c, so LOCALIZE's side condition is satisfied. All that remains is to prove LOCALIZE's two entailments. Let us consider them in reverse order. The second one is $R \vdash \forall x_f.\ \big(\hat{\hat{L}}_2(x_f) \mathbin{-\!*} \hat{\hat{G}}_2(x_f)\big)$, *i.e.*

$$\forall x_f.\ \big(\hat{L}_2(x_f) \mathbin{-\!*} \hat{G}_2(x_f)\big) \vdash \forall x_f.\ \Big(\big((x = x_f) \wedge \hat{L}_2(x_f)\big) \mathbin{-\!*} \big((x = x_f) \wedge \hat{G}_2(x_f)\big) \Big)$$

This turns out to be just a long-winded tautology, and can be done automatically in a tool.

The first of LOCALIZE's entailments is $G_1 \vdash L_1 \mathbin{-\!*} R$, *i.e.*

$$(x=5 \wedge A) \vdash (x=5 \wedge B) * \big(\forall x_f.\ (x_f = 6 \wedge C) \mathbin{-\!*} (x_f = 6 \wedge D)\big)$$

This can be broken into the "variable-related" part $x = 5 \vdash (x = 5) * \big(\forall x_f.\ (x_f = 6 \mathbin{-\!*} x_f = 6)\big)$, which is also a tautology, and the "spatial" part $A \vdash B * (C \mathbin{-\!*} D)$, which was true by assumption above.

With carefully engineering, all of the modified-variable dance above can be done fully automatically, and in a way that is completely hidden to end-users. The only remaining proof goal is the spatial part, which captures the key action of the localization block. To solve these in practice, we end up applying generic lemmas from §5.

---

[3]There is another problem: in the standard model for local variable treatment in separation logic, the separating implication is vacuously true since x cannot simultaneously be both 5 and 6. But since two fatal problems are overkill, let us move on.

```
1  struct Node {int  _Alignas(16) m;
2    struct Node * _Alignas(8) l, * r };
3  void mark(struct Node * x){ //{graph(x, γ)}
4    struct Node * l, * r; int root_mark;
5    if (x == 0) return;
6  //{graph(x, γ) ∧ ∃m, l, r.γ(x) = (m, l, r)}
7  //{graph(x, γ) ∧ γ(x) = (m, l, r)}
8  //↘ {x ↦ m, −, l, r}
9        root_mark = x -> m;
10 //↗ {x ↦ m, −, l, r ∧ m = root_mark}
11 //{graph(x, γ) ∧ γ(x) = (m, l, r) ∧ m = root_mark}
12   if (root_mark == 1) return;
13 //{graph(x, γ) ∧ γ(x) = (0, l, r)}
14 //↘ {x ↦ 0, −, l, r ∧ γ(x) = (0, l, r)}
15 ↯(5.3) l = x -> l; r = x -> r; x -> m = 1;
```

$$16\ //\swarrow \left\{ \begin{array}{l} x \mapsto 1, -, l, r \wedge \gamma(x) = (0, l, r) \wedge \\ \exists \gamma'.\mathit{mark1}(\gamma, x, \gamma') \end{array} \right\}$$

17 $//\{\exists \gamma'.\mathrm{graph}(x, \gamma') \wedge \gamma(x) = (0, l, r) \wedge \mathit{mark1}(\gamma, x, \gamma')\}$

18 $//\{\mathrm{graph}(x, \gamma') \wedge \gamma(x) = (0, l, r) \wedge \mathit{mark1}(\gamma, x, \gamma')\}$

19 $//\searrow \{\mathrm{graph}(l, \gamma')\}$

20 $↯(5.3)$ mark(l);

21 $//\swarrow \{\exists \gamma''.\mathrm{graph}(l, \gamma'') \wedge \mathit{mark}(\gamma', l, \gamma'')\}$

$$22\ // \left\{ \begin{array}{l} \exists \gamma''.\mathrm{graph}(x, \gamma'') \wedge \gamma(x) = (0, l, r) \wedge \\ \mathit{mark1}(\gamma, x, \gamma') \wedge \mathit{mark}(\gamma', l, \gamma'') \end{array} \right\}$$

$$23\ // \left\{ \begin{array}{l} \mathrm{graph}(x, \gamma'') \wedge \gamma(x) = (0, l, r) \wedge \\ \mathit{mark1}(\gamma, x, \gamma') \wedge \mathit{mark}(\gamma', l, \gamma'') \end{array} \right\}$$

24 $//\searrow \{\mathrm{graph}(r, \gamma'')\}$

25 $↯(5.3)$ mark(r);

26 $//\swarrow \{\exists \gamma'''.\mathrm{graph}(r, \gamma''') \wedge \mathit{mark}(\gamma'', r, \gamma''')\}$

$$27\ // \left\{ \begin{array}{l} \exists \gamma'''.\mathrm{graph}(x, \gamma''') \wedge \gamma(x) = (0, l, r) \wedge \\ \mathit{mark1}(\gamma, x, \gamma') \wedge \mathit{mark}(\gamma', l, \gamma'') \wedge \mathit{mark}(\gamma'', r, \gamma''') \end{array} \right\}$$

28 } $//\{\exists \gamma'''.\mathrm{graph}(x, \gamma''') \wedge \mathit{mark}(\gamma, x, \gamma''')\}$

$$
\begin{aligned}
\mathrm{graph}(x, \gamma) \Leftrightarrow & (x = 0 \wedge \mathrm{emp}) \vee \\
& \exists m, l, r.\ \gamma(x) = (m, l, r) \wedge x \bmod 16 = 0 \wedge \quad (4)\\
& x \mapsto m, -, l, r \uplus \mathrm{graph}(l, \gamma) \uplus \mathrm{graph}(r, \gamma)
\end{aligned}
$$

$$\mathit{mark1}(\gamma, x, \gamma') \overset{\Delta}{=} \forall v.\gamma'(v) = \begin{cases} (1, l, r) & \text{when } x = v \wedge \\ & \gamma(v) = (0, l, r) \\ \gamma(v) & \text{otherwise} \end{cases}$$

$$v_1 \overset{\gamma}{\rightsquigarrow}_0 v_2 \overset{\Delta}{=} \exists l, r.\ \gamma(v_1) = (0, l, r) \wedge v_2 \in \{l, r\}$$

$$v_1 \overset{\gamma}{\rightsquigarrow}_0^\star v_2 \overset{\Delta}{=} \text{reflexive, transitive closure of } \text{``} \overset{\gamma}{\rightsquigarrow}_0 \text{''}$$

$$\mathit{mark}(\gamma, x, \gamma') \overset{\Delta}{=} \forall v.\gamma'(v) = \begin{cases} (1, l, r) & \text{when } x \overset{\gamma}{\rightsquigarrow}_0^\star v \wedge \\ & \gamma(v) = (-, l, r) \\ \gamma(v) & \text{otherwise} \end{cases}$$

Fig. 3. Clight code and proof sketch for bigraph mark.

Discussion. The delicacy and detail in the dance above may seem to be making mountains out of molehills, since a careful treatment of modified program variables is hardly a sexy topic. Indeed, in pen-and-paper systems they are molehills, with any number of workarounds including: making local program transformations to introduce fresh variables and arguing for program equivalence, using variables-as-resource [Bornat et al. 2006], or even just sweeping the issue under the rug.

In a mechanized context, working with existing toolsets, these kinds of solutions are not viable. Either we must reinvent a **very, very** large wheel—combined, VST and CompCert contain about 840k LOC—or we must dance within their constraints. VST does not use variables as resource, nor does it have modules to reason about program equivalence. Moreover, it is hardly unique in these respects: most other mechanized verification systems do not support these solutions either [Beckert et al. 2007; Bengtson et al. 2012; Chin et al. 2010; Distefano and Parkinson 2008]. By respecting the design decisions taken by most existing tools, our solutions can be incorporated more easily; in §7 we will see that our additions to VST are less than 1% of its codebase.

3.3 Linked existentials

We have already seen that allowing existentials in postconditions lets us handle modified program variables properly. However, these "linked existentials"—recall that our previous technique hinged on the fact that the existential witness to the variable x_f in the local postcondition L_2 was carried over to the corresponding existential witness in the global postcondition—have other uses as well. To illustrate them, and to demonstrate other aspects of our system, in particular our ability to explore a graph recursively via fold/unfold, we consider another example.

In Figure 3 we put the code and proof sketch of the classic `mark` algorithm that visits and colors every reachable node in a heap-represented graph. The `mark` example contrasts from `find` in several respects. First, it modifies the labels of nodes instead of the edges. Second, each node has two outgoing edges rather than one, so the graph can have a more complex shape. Third, the graph can be nontrivially cyclic. Lastly, the specification we certify (lines 3 and 28) is *local* rather than *global*:

$$\{\mathsf{m\_graph}(x, \gamma)\} \; \mathtt{mark(x)} \; \{\exists \gamma'. \; \mathsf{m\_graph}(x, \gamma') \wedge mark(\gamma, x, \gamma')\}$$

The specification is again stated with mathematical γ, although in this case $\gamma(x)$ maps to triples (m, l, r), where m is a "mark" bit (0 or 1) and $\{l, r\} \subseteq V(\gamma) \uplus \{\mathtt{null}\}$ are the neighbors of v. By "local", we mean that the predicate $\mathsf{m\_graph}(x, \gamma)$ says that the heap represents *only the nodes in γ that are reachable from* x. Rather than passing the entire graph around as `find` does, `mark` will use the fold/unfold relationship given in equation (4), located just under the code in Figure 3, to "unfold" the graph as if it were an inductive predicate.

This fold/unfold relationship deserves attention. First, (4) uses the "overlapping conjunction" $\uplus$ of separation logic; informally $P \uplus Q$ means that P and Q may overlap in the heap (*e.g.*, nodes in the left subgraph can also be in the right subgraph or even be the root x). The presence of the unspecified sharing indicated by the $\uplus$ connective[4] is part of why graph-manipulating algorithms are so hard to verify (*e.g.*, it is hard to apply the FRAME rule). Second, (4) illustrates how industrial-strength settings complicate verification. Lines 1–2 define the data type `Node` used by `mark`. The `_Alignas(n)` directives tell CompCert to align fields on n-byte boundaries. As explained in §5.2, this alignment is necessary in C-like memory models to prove fold-unfold (4), which is why (4) includes an alignment restriction $x \bmod 16 = 0$ and an existentially-quantified "blank" second field for the root $x \mapsto m, -, l, r$. (In our Floyd proofs the alignment restriction and blank second field are nicely hidden "behind the scenes".)

Just as with `find`, the postcondition of `mark` is specified *relationally*, i.e. $\{\exists \gamma'. \; \mathsf{graph}(x, \gamma') \wedge mark(\gamma, x, \gamma')\}$ instead of *functionally*, i.e. $\{\mathsf{graph}(x, mark(\gamma, x))\}$. In the first case *mark* is a relation that specifies that γ' is the result of correctly marking γ from x, whereas in the second *mark* is a function that **computes** a new graph, which is the result of marking γ from x. A relational approach is better for both theoretical and practical reasons. Theoretically, relations are preferable because they are more general. For example, relations allow "inputs" to have no "outputs" (*i.e.* be partial) or alternatively have many outputs (*i.e.* be nondeterministic). Nondeterminism can be quite useful when specifying programs; for example, the CertiCoq garbage collector (§6) is specified nondeterministically to avoid, among other things, specifying how `malloc` allocates fresh blocks of memory. Relations are also preferable to functions because they are more compositional.

Practically, it is painful to define computational functions over graphs in a proof assistant like Coq. For example, Coq requires that all functions terminate, a nontrivial proof obligation over cyclic structures like graphs, but our verification of `mark` is only for partial correctness. Defining relations is much easier because *e.g.* one can use quantifiers and does not have to prove termination. The *mark* and *mark1* relations we use are defined straightforwardly at the bottom of Figure 3.

The highlights of the proof are as follows. In lines 8–10, imagine unfolding the graph predicate in line 7 using equation (4) and then zooming in to the root node x for lines 8–10, before zooming back out in line 11. Lines 18–22 of Figure 3 contain an example where we use the power of linked existentials in the LOCALIZE rule to "extract" the existentially-quantified γ'' from inside the localization block to outside it. The rest of the proof is relatively routine.

---

[4]Recall that the standard semantics of the separation logic connectives used in this paper are in Figure 7 on page 13.

3.4 Additional verified examples

In addition to the proof of `find` shown above, we do a proof of `union` in the same style. We also do a *second* version of union-find, this time using arrays rather than `malloc`-allocated nodes. The mathematical definitions, *e.g.* for *findS*, are entirely shared, demonstrating that we have separated the concerns of abstract algorithmic reasoning from the nitty-gritty details of heap representation.

We also have done a version of `mark` for DAGs (acyclic). In general, using DAGS are both easier and harder than cyclic graphs. On the one hand, we get genuine separation between the root and its children; on the other hand, we need to maintain acyclicity if we modify the link structure. As a more aggressive example of modifying the link structure of a graph, we have verified `spanning`, which prunes a graph into its spanning tree; for space reasons we put this example in Appendix A. Our flagship example, the CertiCoq garbage collector, will be discussed in §6.

4 A REUSABLE LIBRARY OF FORMALIZED GRAPH THEORY

To verify the functional correctness of graph algorithms it is natural to want to use graph theory to describe program behavior. As discussed in §8, 25 years of research into mechanized graph theory can be succinctly summarized as "it is a little tricky". Accordingly, an essential component of our work is a library of formalized graph theory that is actually powerful enough, and expressive enough, to mechanically verify realistic algorithms written in C. As will be shown in §7, our mathematical graph constructions comprise a considerable fraction of our codebase, so it is vital that our framework be highly modular to enable reuse of definitions and proofs from one example to the next. We present our mathematical graph framework with an emphasis on this modularity.

4.1 Definitions of graphs
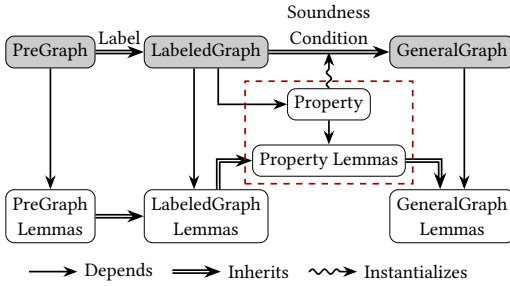
Fig. 4. Structure of the Mathematical Graph Library

Fig. 5. A PreGraph with valid/invalid vertices and edges.

Our first challenge is that graph theory is usually based on *set theory* but our formalization in Coq is based on *type theory*. We choose to formalize graph theory directly in Coq instead of formalizing set theory and then building graph theory atop it to take advantage of Coq's built-in support for type-related constructions. To balance the dichotomy between the generality and the speciality of the library, we divide the concept of graph into three structures, PreGraph, LabeledGraph and GeneralGraph, arranged in a hierarchy. Figure 4 shows the architecture of the library.

PreGraph. A PreGraph is a hextuple $(\mathcal{V}, \mathcal{E}, \mathsf{V}, \mathsf{E}, \mathsf{src}, \mathsf{dst})$. $\mathcal{V}$ and $\mathcal{E}$ are the underlying carrier types of vertices and edges. V and E are predicates over $\mathcal{V}$ and $\mathcal{E}$ that specify the notion of *validity* in the graph. Finally, src and dst map each edge to its source and destination.

The benefits of introducing validity are twofold. The first is a neat resolution of the incompatibility between type theory and set theory. In set theory, one element can belong to multiple sets, and adding or removing vertices or edges is as easy as altering the set directly to represent the result of the operation. In type theory, however, a term can only belong to one type, which makes it difficult

$$\text{path} \triangleq (v_0, [e_0, e_1, \ldots, e_k])$$

$$\text{s\_evalid}(\gamma, e) \triangleq \mathsf{E}_\gamma(e) \wedge \mathsf{V}_\gamma(\mathsf{src}_\gamma(e)) \wedge \mathsf{V}_\gamma(\mathsf{dst}_\gamma(e))$$

$$\text{valid\_path}\big(\gamma, (v, [])\big) \triangleq \mathsf{V}_\gamma(v)$$

$$\text{valid\_path}\big(\gamma, (v, [e_1, e_2, \ldots, e_n])\big) \triangleq v = \mathsf{src}_\gamma(e_1) \wedge \text{s\_evalid}(\gamma, e_1) \wedge$$
$$\mathsf{dst}_\gamma(e_1) = \mathsf{src}_\gamma(e_2) \wedge$$
$$\text{s\_evalid}(\gamma, e_2) \wedge \ldots \wedge \mathsf{dst}_\gamma(e_{n-1}) = \mathsf{src}_\gamma(e_n)$$

$$\text{end}\big(\gamma, (v, [])\big) \triangleq v$$

$$\text{end}\big(\gamma, (v, [e_1, e_2, \ldots, e_n])\big) \triangleq \mathsf{dst}_\gamma(e_n)$$

$$\gamma \vDash s \overset{p}{\rightsquigarrow} t \triangleq \text{valid\_path}(\gamma, p) \wedge \text{fst}(p) = s \wedge \text{end}(\gamma, p) = t$$

$$\gamma_1 \cong \gamma_2 \triangleq \forall e.\, \mathsf{E}_{\gamma_1}(e) \Leftrightarrow \mathsf{E}_{\gamma_2}(e) \wedge \forall v.\, \mathsf{V}_{\gamma_1}(v) \Leftrightarrow \mathsf{V}_{\gamma_1}(v) \wedge$$
$$\forall e.\, \mathsf{E}_{\gamma_1}(e) \Rightarrow \mathsf{src}_{\gamma_1}(e) = \mathsf{src}_{\gamma_2}(e) \wedge \mathsf{dst}_{\gamma_1}(e) = \mathsf{dst}_{\gamma_2}(e)$$

$$\gamma \smallsetminus S \triangleq \mathcal{V}_{\gamma'} = \mathcal{V}_\gamma \wedge \mathcal{E}_{\gamma'} = \mathcal{E}_\gamma \wedge$$
$$\mathsf{src}_{\gamma'} = \mathsf{src}_\gamma \wedge \mathsf{dst}_{\gamma'} = \mathsf{dst}_\gamma \wedge$$
$$\gamma'_\mathsf{V} = \Big(\lambda x.\, \gamma_\mathsf{V}(x) \wedge \neg S_\mathsf{V}(x)\Big) \wedge$$
$$\gamma'_\mathsf{E} = \Big(\lambda x.\, \gamma_\mathsf{E}(x) \wedge \neg S_\mathsf{E}(x)\Big)$$

$$\text{MathGraph}(\gamma) \triangleq \Big\{ \text{null} : V$$
$$\text{valid\_graph} : \forall e.\, \text{evalid}(\gamma, e) \Rightarrow \text{vvalid}\big(\gamma, \mathsf{src}(\gamma, e)\big) \wedge$$
$$\big(e = \text{null} \vee \text{vvalid}(\gamma, e)\big)$$
$$\text{valid\_not\_null} : \forall v.\, \text{vvalid}(\gamma, v) \Rightarrow v \neq \text{null}\Big\}$$

$$\text{LstGraph}(\gamma) \triangleq \Big\{ \text{out} : V \to E$$
$$\text{only\_one\_edge} : \forall v, e.\, \text{vvalid}(\gamma, v) \Rightarrow$$
$$\Big(\mathsf{src}(\gamma, e) = v \wedge \text{evalid}(\gamma, e)\Big) \Leftrightarrow e = \text{out}(v)$$
$$\text{acyclic\_path} : \forall v, p.\, \gamma \vDash v \overset{p}{\rightsquigarrow} v \Rightarrow p = (v, [])\Big\}$$

$$\text{FiniteGraph}(\gamma) \triangleq \Big\{ \text{finite\_v} : \exists S_v, M_v \text{ s.t. } |S_v| \le M_v \wedge \forall v.\, \text{vvalid}(\gamma, v) \Rightarrow v \in S_v$$
$$\text{finite\_e} : \exists S_e, M_e \text{ s.t. } |S_e| \le M_e \wedge \forall e.\, \text{evalid}(\gamma, e) \Rightarrow e \in S_e\Big\}$$

Fig. 6. Some Graph definitions

to analogously change the type to represent the result. As is common practice, the predicates V and E specify whether a vertex/edge is *valid* (in the graph) or *invalid* (out). Adding or removing vertices/edges is as simple as weakening or strengthening these two predicates.

The second benefit is the ability to represent incomplete graphs. Consider starting from a graph γ and then removing a subgraph; the remaining "doughnut" structure is **not** necessarily a graph, since there may be dangling edges pointing into the "hole". Figure 5 shows just such a situation, where a connected graph (everything is reachable from v_0) has had the connected subgraph reachable from v_1 removed; *e.g.*, the edge e is dangling. The last conjunct in the *findS* relation from Figure 1 is an example of where a real verification needs to reason about just such a doughnut, in particular to specify that the unreachable portion of a graph has not changed.

We define many fundamental graph concepts on PreGraphs, including structures like *path**, predicates such as *is_cyclic* and *reachable**, operations such as *add_vertex* and *remove_edge*, and relations between PreGraphs such as *structurally_identical** and *subgraph*. Definitions of the concepts marked with asterisks are shown in Figure 6 to give a flavor of the subtleties involved in getting definitions that really work. These general concepts, together with around 500 derived lemmas, provide a solid foundation for more specific theorems needed in concrete verifications.

LabeledGraph. A LabeledGraph is septuple (PreGraph, $\mathcal{L}_\mathcal{V}$, $\mathcal{L}_\mathcal{E}$, $\mathcal{L}_\mathcal{G}$, vl, el, gl) that augments a PreGraph with *labels* on vertices, edges, and/or the graph as a whole. $\mathcal{L}_\mathcal{V}$, $\mathcal{L}_\mathcal{E}$, and $\mathcal{L}_\mathcal{G}$ are the associated carrier types; vl, el, and gl are the labeling functions themselves. The need for such labels is that many classic graph problems from union-find (node ranks) to Dijkstra (edge weights) require them. The need for a label on the graph as a whole is a little more subtle; in §6 we use one in the garbage collector to keep track of *e.g.* the number of generations and their boundaries. Since LabeledGraph is built on PreGraph, it inherits all PreGraph lemmas via Coq's *type coercion* mechanism, while enabling additional lemmas involving labels.

GeneralGraph. PreGraphs and LabeledGraphs let us state and prove many useful lemmas that follow essentially by the nature of our graph constructions. However, when proving the correctness of graph algorithms, we often need more specificity in our mathematical graphs so that we may model the real program's behaviors closely. For example, the uf_graph used in find restricted each vertex to having exactly one out-edge. On the other hand, these restrictions vary greatly by algorithm, so we do not want to bake them into our core definitions. We achieve this flexibility using GeneralGraphs, which augment LabeledGraphs by adding arbitrarily complex "soundness conditions", indicated in Figure 4 with a dashed border. Further, the type coercion we described earlier continues to apply, meaning that a GeneralGraph can seamlessly behave like a LabeledGraph or a PreGraph, thereby inheriting their lemmas. This combination of specificity and generality is what makes GeneralGraphs versatile. Moreover, we can compose complicated soundness conditions from reusable pieces, further enabling code sharing between algorithms.

4.2 Composing soundness plugins

Soundness conditions are often specific to each algorithm, but they feature some recurring themes. We take advantage of this pattern by developing *soundness plugins*, *i.e.* definitions of soundness conditions along with related lemmas. By combining these plugins we can describe the soundness condition we need for a particular algorithm. When proving lemmas about the resulting combination, we can use known facts about the separate plugins, in addition to lemmas that emerge due to the various combinations. This complexity is managed smoothly by Coq's typeclass system, increasing the compositionality of the system. Consider the following oft-used graph properties:

- MathGraph*: all valid edges must have a valid source vertex; the destination vertex must either be valid or be a special invalid node called null
- FiniteGraph*: the sets of valid vertices and edges are both finite
- LstGraph*: the graph is list-like: each vertex has only one outgoing edge; no nontrivial loops
- BiGraph: there are exactly two outgoing edges per node

$$\sigma \models P * Q \quad \overset{\Delta}{=} \quad \exists \sigma_1, \sigma_2.\, (\sigma_1 \oplus \sigma_2 = \sigma) \wedge (\sigma_1 \models P) \wedge (\sigma_2 \models 2)$$

$$\sigma \models P \mathbin{-\!\!*} Q \quad \overset{\Delta}{=} \quad \forall \sigma_1, \sigma_2.\, (\sigma_1 \oplus \sigma = \sigma_2) \wedge (\sigma_1 \models P) \Rightarrow (\sigma_2 \models Q)$$

$$\sigma \models P \mathbin{-\!\circledast} Q \quad \overset{\Delta}{=} \quad \exists \sigma_1, \sigma_2.\, (\sigma_1 \oplus \sigma = \sigma_2) \wedge (\sigma_1 \models P) \wedge (\sigma_2 \models Q)$$

$$\sigma \models P \mathbin{\uplus} Q \quad \overset{\Delta}{=} \quad \exists \sigma_1, \sigma_2, \sigma_3.\, (\sigma_1 \oplus \sigma_2 \oplus \sigma_3 = \sigma) \wedge (\sigma_1 \oplus \sigma_2 \models P) \wedge (\sigma_2 \oplus \sigma_3 \models Q)$$

Fig. 7. Separation logic connectives; $\oplus$ is the join operation on states, *e.g.* a disjoint union on heaps

Definitions of the concepts marked with asterisks are shown in Figure 6 for illustration.

We can compose LstGraph, MathGraph, and FiniteGraph together into a new plugin called LiMaFin, which, incidentally, is the soundness condition of mathematical uf_graph we used to verify `find` in Figure 1. In our verification of `mark` in Figure 3, we use a similar soundness condition BiMaFin, which uses BiGraph instead of LstGraph. The commonalities and differences between LiMaFin and BiMaFin are readily apparent from their construction.

5 DEFINING AND REASONING ABOUT SPATIAL GRAPHS

To prove the functional correctness of graph-manipulating algorithms implemented in a real language, we need to connect the heap representation of graphs, the memory model of the programming language, and the mathematical properties of graphs from §4. The first of these turns out to be surprisingly subtle as we shall see in §5.1 and §5.2. The main challenge for the others is to engineer a framework that is generic enough and modular enough to be useful in practice in a variety of settings; we cover it in §5.3.

5.1 Recursive definitions yield poor graph predicates

Recursive predicates are ubiquitous in separation logic—so much so that when one writes the definition of a predicate as $P \overset{\Delta}{=} \ldots P \ldots$, no one raises an eyebrow despite the dangers of circularity in mathematics. Indeed, the vast majority of the time there is no danger thanks to the magic of the Knaster-Tarski fixpoint μ_T [Tarski 1955]. Formally, one does not define P directly, but rather defines a functional $F_P \overset{\Delta}{=} \lambda P.\ \ldots P \ldots$ and then defines P itself as $P \overset{\Delta}{=} \mu_T F_P$. Assuming, as one typically does without comment, that F_P is *covariant*, i.e. $(P \vdash Q) \Rightarrow (F\,P \vdash F\,Q)$, one then enjoys the fixpoint equation $P \Leftrightarrow \ldots P \ldots$, formally justifying the typically-written pseudodefinition ("$\overset{\Delta}{=}$").

Suppose we define a graph predicate graph_T this way, *e.g.* along the lines of the fold/unfold definition in Figure 3:

$$\text{graph}_T(x, \gamma) \overset{\Delta}{=} \Big(x = 0 \wedge \text{emp} \Big) \vee$$
$$\Big(\exists m, l, r.\ \gamma(x) = (m, l, r) \wedge \big(x \mapsto m, l, r \mathbin{\uplus} \text{graph}_T(l, \gamma) \mathbin{\uplus} \text{graph}_T(r, \gamma) \big) \Big)$$

Although we can apply Knaster-Tarski (because the functional needed to define graph_T is covariant), the result is hard to use. Consider the following memory m for a toy machine:

Clearly $m \models 100 \mapsto 42, 100, 0$. But it seems also clear that this memory represents a one-cell cyclic graph as illustrated in the accompanying diagram, *i.e.* we want $m \models \text{graph}_T(100, \hat{\gamma})$, where $\hat{\gamma}(100) = (42, 100, 0)$. This is equivalent to wanting to be able to prove $100 \mapsto 42, 100, 0 \vdash \text{graph}_T(100, \hat{\gamma})$. Unfortunately, as explained in Appendix C, this is rather difficult to do since applying the natural proof techniques actually strengthens the goal. In fact we do not know if this entailment is

address	value
102	0
101	100
100	42

provable, but the difficulties encountered in proving what "should be" straightforward suggest that Knaster-Tarski should be treated with caution when defining spatial predicates for graphs.

The other direction, $\text{graph}_T(100, \hat{\gamma}) \vdash 100 \mapsto 42, 100, 0$, **is** true but is not easy to prove, relying on the constructions in §5.2 and the fact that μ_T constructs the least fixpoint. In contrast, $\text{graph}_T(100, \hat{\gamma}) \vdash 100 \mapsto 42, 100, 0 * \top$ is easy.

5.2 Defining a good graph predicate

Rather than trying to define graph as a recursive fixpoint, we will instead give it a flat structure. Graphs in separation logic have been defined in similar ways before, *e.g.* [Sergey et al. 2015]; our innovation is that we prove—with the amount of precision required to convince Coq—that we can still enjoy fold/unfold with our flat definition. Our path starts with the iterated separating conjunction or "big star", which is first defined over lists and then extended to sets as follows:

$$\underset{\{l_1, l_2, \ldots, l_n\}}{\bigstar} P \stackrel{\Delta}{=} P(l_1) * P(l_2) * \ldots * P(l_n) \quad \Big| \quad \underset{S}{\bigstar} P \stackrel{\Delta}{=} \exists L. \ (\text{NoDup } L) \wedge (\forall x. \ x \text{ in } L \Leftrightarrow x \in S) \wedge \underset{L}{\bigstar} P$$

We are now ready to define a good graph predicate: $\text{graph}(x, \gamma) \stackrel{\Delta}{=} \underset{v \in reach(\gamma, x)}{\bigstar} v \mapsto \gamma(v)$

The graph γ here need not be a bigraph, but *e.g.* can have many edges.

Our definition of graph is flat in the sense that there is no obvious way to follow the link structure recursively. Happily, we can recover a general recursive fold/unfold (if $x \mapsto \gamma(x)$ and the GeneralGraph has the necessary properties in its soundness condition):

$$\text{graph}(x, \gamma) \Leftrightarrow x \mapsto \gamma(x) \uplus \Big(\underset{n \in neighbors(\gamma, x)}{\biguplus} \text{graph}(\gamma, n) \Big) \quad \text{where} \quad \underset{l_1, \ldots, l_n}{\biguplus} P \stackrel{\Delta}{=} P(l_1) \uplus \ldots \uplus P(l_n) \tag{5}$$

The proof of the $\Leftarrow$ direction requires care. The difficulty is that if two nodes $x \mapsto \gamma(x)$ and $x' \mapsto \gamma(x')$ are *skewed*, *i.e.* "partially overlapping" with some—but not all—of x's memory cells shared with x', then the $\bigstar$ on the left hand side cannot separate them. To avoid skewing we require that $x \mapsto \gamma(x)$ be *alignable*. A predicate P is alignable when

$$\forall x, y. \ \Big(P(x) \uplus P(y) \vdash \big(P(x) \wedge x = y \big) \vee \big(P(x) * P(y) \big) \Big)$$

That is, they overlap either completely or not at all. In a Java-like memory model this property is automatic because pointers in such a model always point to the root/beginning of an object. In contrast, in a C-like memory model such as in VST/CompCert, this property is not automatic because pointers can point anywhere. In such a model, alignment is most easily enforced by storing graph nodes at addresses that are multiples of an appropriate size (16 in Figure 3).

Some of our VST proofs, *e.g.* for `find`, do not use fold/unfold, instead preferring to use the lemmas in §5.3 directly. Others, *e.g.* `mark`, do—and it is helpful to have both options avilable. We also prove fold/unfold lemma for DAGs in which we get a $*$ between the root and its $\uplus$-joined neighbors, rather than the $\uplus$ present in (5).

5.3 Ramification Libraries

VST employs a somewhat unusual Step-Indexed heap model to represent spatial predicates, and uses it to support an unusually rich program logic. [Appel et al. 2014] However, the spatial representation of our graph library does not rely on any of its bells and whistles. To isolate our development from these unnecessary complications, we use two interfaces: Core Logic and Supplementary Logic, where Supplementary Logic is built upon the much simpler and more universal Direct Model of heap representation. We present the architecture of our spatial development in Figure 8, where it should be straighforward to see that both models can instantiate both interfaces.

687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
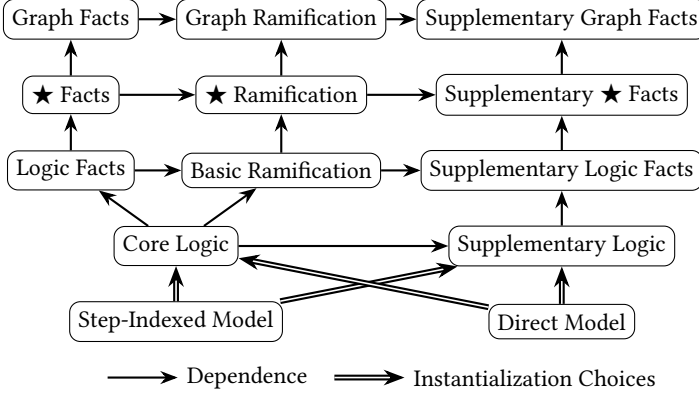
Fig. 8. Infrastructure of ramification library

Generally speaking, our VST proofs only need the Core properties to prove our examples. Each interface defines some operators of separation logic and provides some axioms about how they work. For example, $*$ and $-\!\!*$ are in Core Logic, along with the axiom $(P \vdash Q -\!\!* R) \Leftrightarrow (P * Q \vdash R)$. On the other hand, the $\uplus$ and $-\!\!\circledast$ operators are in Supplementary Logic, along with rules like $P \vdash P \uplus P$.

Above the Logic layer we have three towers, each three levels high. The tower on the left contains basic lemmas about Logic, $\bigstar$, and graph. For instance, in the $\bigstar$ Facts box we prove:

$$\frac{A \cap B = \emptyset}{\underset{x \in A}{\bigstar} P(x) \; * \; \underset{x \in B}{\bigstar} P(x) \;\; \Leftrightarrow \;\; \underset{x \in A \cup B}{\bigstar} P(x)}$$

The middle tower is more interesting in that it is entirely focused on ramification entailments. A robust library of ramification entailments is essential to make ramification work smoothly in practice. The Basic Ramification box in the lower layer contains lemmas like:

$$\frac{G_1 \vdash L_1 * \forall x. (L_2 -\!\!* G_2) \qquad G_1' \vdash L_1' * \forall x. (L_2' -\!\!* G_2')}{G_1 * G_1' \vdash (L_1 * L_1') * \forall x. \big((L_2 * L_2') -\!\!* (G_2 * G_2')\big)}$$

We use this lemma to break large ramification entailments into more manageable pieces in a compositional way. The $\bigstar$ Ramification box in the middle layer contains lemmas like:

$$\frac{A \cap B = \emptyset \qquad A' \cap B = \emptyset}{\underset{x \in A \cup B}{\bigstar} P(x) \vdash \underset{x \in A}{\bigstar} P(x) * \left(\underset{x \in A'}{\bigstar} P(x) -\!\!* \underset{x \in A' \cup B}{\bigstar} P(x) \right)}$$

The Graph Ramification box in the top layer is focused lemmas such as the following "update one node" lemma, which was used on line 15 in Figure 3:

$$\frac{\forall x_0 \neq x. \; \gamma(x_0) = \gamma'(x_0) \qquad \mathsf{neighbors}(\gamma, x) = \mathsf{neighbors}(\gamma', x)}{\mathsf{graph}(x, \gamma) \vdash x \mapsto \gamma(x) * \big(x \mapsto \gamma'(x) -\!\!* \mathsf{graph}(x, \gamma')\big)}$$

This layered structure enables proof reuse. All of the theorems for graph are proved from the properties of iterated separating conjunction, but having a modular library allows $\bigstar$ to be reused in other structures smoothly. Further, all our verifications of different graph algorithms use the

proof rules of graph at the top level in the library. Taking the marking algorithm we introduced in §2 as an example, we prove the following theorem from the library:

$$\frac{n \in \text{neighbors}(\gamma, x)}{\text{graph}(x, \gamma) \vdash \text{graph}(n, \gamma) * (\forall \gamma'. \, mark(\gamma, n, \gamma') \wedge \text{graph}(n, \gamma') \mathbin{-\!\!*} mark(\gamma, n, \gamma') \wedge \text{graph}(x, \gamma'))}$$

The Supplementary tower contains properties not used by most of the VST examples. This includes the fold/unfold relationship from §5.2, facts about precision, *etc.* These are currently included mostly for completeness, but do make our library more general should we wish to accommodate an alternate prover that uses the Direct Model.

6 CERTIFYING A GARBAGE COLLECTOR FOR CERTICOQ

6.1 Background

The CertiCoq compiler [Anand et al. 2017] translates Gallina code to Clight, which CompCert compiles to assembly [Leroy 2006]. CertiCoq's garbage collector (GC) is written in Clight and supports Gallina's assumption of infinite memory. CertiCoq aims to be end-to-end certified, so the GC must be too.

The 12-generation collector, written in the spirit of the OCaml GC, is relatively realistic and sophisticated, though by no means industrial-strength. Because CertiCoq borrows OCaml's representation of blocks and values [Hickey et al. 2014], the GC must support features such as variable-length memory objects, object fields that may be boxed or unboxed and must be disambiguated at runtime, pointers to places outside the GC-managed heap, etc. The CertiCoq GC's task is a little easier than the OCaml GC's because its mutator is purely functional[5]. The mutator maintains an args array of local variables, which the GC scans to calculate the root set. When called, the GC collects the first generation into the second using Cheney's algorithm [Cheney 1970]. This collection may trigger the collection of the second generation into the third, etc., and the GC completes this potential cascade before returning control to the mutator. A fuller explanation of GC's operation is in Appendix D.

The mutator's args array of local variables is critical in the GC's specification. The GC must ensure that all memory objects that the mutator can reach by recursively following the fields of args *before* the collection can still be reached, via the same steps, *after* the collection. This problem can be abstracted into mathematical graphs, where we must prove graph isomorphism.

6.2 From C to Mathematical Graphs

In the code, the metainformation of the 12 generations is stored as an 12-element array heap. Each generation, as a memory segment, is represented as a struct space, which contains three pointers: start marking the start address of a generation, next representing the next available address in a generation, and limit marking the last address of a generation. The basic unit manipulated by the garbage collector is a chunk of memory called a block. Blocks can be of different sizes; the size of a particular block is stored in its header (using 22 bits of the word stored at offset $-sizeof(\text{void}*)$), and the remainder of the block is a continuous array of fields. Each field is either an unboxed integer data value, or a pointer, which may be either within the GC or to external structures. To disambiguate the two, we follow OCaml's practice of requiring that all pointers are even-aligned and that all integers to be odd (essentially, to be only 31 bits long) [Hickey et al. 2014].

From the perspective of the algorithm, the 12 generations can be seen as a graph γ. Each block can be can be seen as a vertex and pointers to other blocks indicate edges between vertices. More formally, we decide to encode each vertex of this graph as a pair of natural numbers (v_g, v_i) which means the vertex is the v_ith block in the v_gth generation. We encode each edge as a pair of vertex

---

[5]That is: Gallina is purely functional, and the Clight code generated by CertiCoq preserves this behavior.

and index (v, i) which means this edge is from vertex v and the associated pointer is in the ith field of the corresponding block. The source function always satisfies $src(\gamma, (v, i)) = v$. Each vertex is labeled with the integer data items and the indices of pointers in the fields. There is also a global label of γ which has the start/limit addresses and number of vertices of each generation. We can reconstruct the 12 generations in memory from γ under this setting without redundancy. For example, to determine the `next` pointer of a generation, we can sum the sizes of each vertex in that generation using its label, and then add the `start` address.

6.3 Forward

The function `forward` is the GC's workhorse. When correctly given the spaces `from` and `to` and a pointer p to a memory block in `from`, it copies the memory block to the next available location in the `to` space. The function is robust: if passed a "pointer" argument that is actually a data value, or that points outside of `from`, it behaves appropriately by taking no action. As we will see in §6.4, these checks are nontrivial. The function is also versatile: it is used to collect the mutator's `args` (which are ∗-separated from the heap) and also to collect the blocks *in* the heap that are reachable via `args`. Its behavior needs to be subtly different in these two cases. Figure 9 shows a decorated proof sketch of `forward` in the latter case, which is harder to verify.

Two abstractions of `struct thread_inf` and `file_info`—*finf* and *tinf*—together allow us to extract the mutator's `args` array, and The proposition *super_compatible* encapsulates various checks about *e.g.* legal bounds that avoid some overflow issues. For concision, the facts known to us in *e.g.* line 1 are represented by ϕ_1, and then ϕ_1 can feature as a fact *e.g.* in line 8.

Line 16 shows the case when the block passed to the function was already forwarded. This may seem strange, but is vital because the same block may be reachable from the `args` via different paths. Such a block's header is zeroed out and its 0th field holds the address of its copy[6], so we simply reroute to that copy. Line 18 shows that this operation gives us a new graph, $\gamma' = upd\_edge(\gamma, e, copy(\gamma, v'))$. This means to reroute the edge e in γ and make it point at $copy(\gamma, v')$.

Lines 24 to 26 show how a block is copied over to the next-available spot in the `to` space. Some of the grungy details having to do with variable-sized memory blocks being to show up in the C code, but the verification looks relatively clean thanks to our mathematical graph framework. This is just the copying of a vertex, and so our new graph after the change is $\gamma' = copy\_vertex(\gamma, to, v', v'')$. The final step (line 32) is to reroute to this new copy, and this is done just as in line 16. The resultant graph is $\gamma'' = upd\_edge(\gamma', e, v'')$.

The postcondition is a little different from those of `find` and `mark` seen earlier, in that it does not provide a relation saying that `forward` has been functionally correct. Rather, we defined the relation to reflect the result of operations in `forward`, such as a vertex is copied, an edge is redirected, and etc,. For a taste, we put the complete definition of *forward_relation* in Appendix E. We have relations of this sort for all the key functions in our GC, and we show that the composition of these correctly corresponds to the C code of our garbage collector. The final functional correctness is derived from these relations.

6.4 Performance and Overflows and Undefined Behaviours, Oh My!

Bugs in the GC code. We discovered and fixed two bugs in the source code during our verification. The first was a performance bug we discovered when developing the key invariants. The original GC code executed Cheney's algorithm too conservatively, scanning the entire `to` space for backward pointers into `from`. We showed that scanning a subset of `to` suffices. Performance doubled.

---

[6]These guarantees are set up by `forward` itself. Refer to lines 28 and 29 of Figure 9 to see this being done.

Figure content (Fig. 9), left column:

$$1 \quad /\!\!/ \left\{ \begin{array}{l} \forall \gamma, \textit{finf}, \textit{tinf}, \textit{from}, \textit{to}, v, n. \\ \textrm{gc\_graph}(\gamma) * \textrm{finf}(\textit{finf}) * \textrm{tinf}(\textit{tinf}) \wedge \\ \textit{compat}(\gamma, \textit{finf}, \textit{tinf}, \textit{from}, \textit{to}) \wedge \\ \textrm{s} = \textit{start}(\gamma, \textit{from}) \wedge \textrm{l} = \textrm{s} + \textit{gensize}(\gamma, \textit{from}) \\ \wedge \; n = \textit{naddr}(\textit{tinf}, \textit{to}) \wedge \textrm{p} = \textit{vaddr}(\gamma, v) + n \end{array} \right\} \stackrel{\Delta}{=} \phi_1$$

```
2  void forward (value *s, *l, **n, *p,
3                            int depth) {
4   value * v; value va = *p;
5   if(Is_block(va)) {//is ptr
6     v = (value*)((void *)va);
7     if(Is_from(s, l, v)) {//in from
```

$$8 \quad /\!\!/ \left\{ \begin{array}{l} \phi_1 \wedge \exists e, v'. \; \textit{lab}(\gamma, v)[n] = e \wedge \\ \textit{dst}(\gamma, e) = v' \wedge v = \textit{vaddr}(\gamma, v') \end{array} \right\} \stackrel{\Delta}{=} \phi_8$$

$$9 \quad /\!\!/ \searrow \left\{ \begin{array}{l} \exists \textit{flds}', \textit{hdr}'. \; \textit{flds}' = \textit{lab}(\gamma, v') \wedge \\ v' \mapsto \textit{flds}' \wedge \textit{hdr}' = \textit{flds}'[-1] \end{array} \right\} \stackrel{\Delta}{=} \phi_9$$

```
10    header_t hd = Hd_val(v);
```

$$11 \quad /\!\!/ \nearrow \left\{ \phi_9 \wedge \textrm{hd} = \textit{val}(\textit{hdr}') \right\}$$

$$12 \quad /\!\!/ \left\{ \phi_8 \wedge \textrm{hd} = \textit{val}(\textit{hdr}') \right\}$$

```
13      if(hd == 0) {//already forwarded
```

$$14 \quad /\!\!/ \left\{ \phi_8 \wedge \textrm{hd} = 0 \right\} \stackrel{\Delta}{=} \phi_{14}$$

$$15 \quad /\!\!/ \searrow \left\{ \begin{array}{l} \exists \textit{flds}, \textit{flds}'. \; v \mapsto \textit{flds} \wedge v' \mapsto \textit{flds}' \wedge \\ \textit{flds} = \textit{lab}(\gamma, v) \wedge \textit{flds}' = \textit{lab}(\gamma, v') \wedge \\ \textit{flds}'[0] = \textit{vaddr}(\gamma, \textit{copy}(\gamma, v')) \wedge \\ \textrm{p} = \&\textit{flds}[n] \end{array} \right\} \stackrel{\Delta}{=} \phi_{15}$$

```
16      *p = Field(v,0);
```

$$17 \quad /\!\!/ \nearrow \left\{ \phi_{15} \wedge \textit{flds}[n] := \textit{flds}'[0] \right\}$$

Figure content (Fig. 9), right column:

$$18 \quad /\!\!/ \left\{ \begin{array}{l} \phi_{31} \wedge \exists \gamma'. \; \textrm{gc\_graph}(\gamma') \wedge \\ \gamma' = \textit{upd\_edge}(\gamma, e, \textit{copy}(\gamma, v')) \wedge \\ \textit{postcondition}(\gamma, \gamma', \textit{tinf}, \textit{finf}, \textit{from}, \textit{to}) \end{array} \right\}$$

```
19    } else {//not yet forwarded
```

$$20 \quad /\!\!/ \left\{ \phi_8 \wedge \textrm{hd} \neq 0 \right\} \stackrel{\Delta}{=} \phi_{20}$$

```
21      int i; int sz; value *new;
22      sz = size(hd); new = *n+1; *n = new+sz;
```

$$23 \quad /\!\!/ \left\{ \begin{array}{l} \phi_{20} \wedge \textrm{sz} = \textit{blocksize}(\textrm{hd}) \wedge \\ \textrm{new} = \textit{start}(\gamma, \textit{to}) + \textit{used}(\gamma, \textit{to}) + 1 \wedge \textrm{n} = \textrm{new} + \textrm{sz} \end{array} \right\} \stackrel{\Delta}{=} \phi_{23}$$

```
24      Hd_val(new) = hd;
25      for(i = 0; i < sz; i++)
26        Field(new, i) = Field(v, i);
```

$$27 \quad /\!\!/ \left\{ \begin{array}{l} \phi_{23} \wedge \exists \gamma', v'', \textit{tinf}'. \; \textrm{gc\_graph}(\gamma') \wedge \\ v'' = \textit{newly\_copied\_vertex}(\gamma, \textit{to}) \wedge \\ \gamma' = \textit{copy\_vertex}(\gamma, \textit{to}, v', v'') \wedge \\ \textit{compat}(\gamma', \textit{finf}, \textit{tinf}', \textit{from}, \textit{to}) \end{array} \right\} \stackrel{\Delta}{=} \phi_{27}$$

```
28      Hd_val(v) = 0;
29      Field(v, 0) = (value)((void *)new);
```

$$30 \quad /\!\!/ \left\{ \phi_{27} \wedge \textit{val}(\textit{hdr}') = 0 \wedge \textit{flds}'[0] = \textit{copy}(\gamma, v') \right\} \stackrel{\Delta}{=} \phi_{30}$$

$$31 \quad /\!\!/ \searrow \left\{ \exists \textit{flds}. \; v \mapsto \textit{flds} \wedge \textit{flds} = \textit{lab}(\gamma', v) \right\} \stackrel{\Delta}{=} \phi_{31}$$

```
32      *p = (value)((void *)new);
```

$$33 \quad /\!\!/ \nearrow \left\{ \phi_{31} \wedge \textit{flds}[0] := \textit{vaddr}(\gamma, v'') \right\}$$

$$34 \quad /\!\!/ \left\{ \phi_{30} \wedge \exists \gamma''. \; \textrm{gc\_graph}(\gamma'') \wedge \gamma'' = \textit{upd\_edge}(\gamma', e, v'') \right\}$$

```
35  }}}}
```

$$36 \quad /\!\!/ \left\{ \textit{postcondition}(\gamma', \gamma'', \textit{tinf}', \textit{finf}, \textit{from}, \textit{to}) \right\}$$

$$\textit{postcondition}(\gamma, \gamma', \textit{tinf}', \textit{finf}', \textit{from}, \textit{to}) \stackrel{\Delta}{=} \textrm{gc\_graph}(\gamma') * \textrm{finf}(\textit{finf}') * \textrm{tinf}(\textit{tinf}') \wedge$$
$$\textit{compat}(\gamma', \textit{finf}', \textit{tinf}', \textit{from}, \textit{to}) \wedge \textit{forward\_relation}(\gamma, \gamma', \textit{from}, \textit{to})$$

Fig. 9. Clight code and proof sketch for forward

The second bug was an overflow when subtracting two pointers to calculate the size of a space, as below. Pointers `start` and `limit` point to the beginning and end of the $i^{\textrm{th}}$ space of the heap `h`.

```
int w = h->spaces[i].limit - h->spaces[i].start;
```

This subtraction is defined in C and Clight, but will overflow if the difference equals or exceeds 2^{31}. We adjusted the size of the largest generation to avoid this overflow.

Undefined behavior in C. We found two places where the semantics of Clight was unable to specify an OCaml-style GC. The first involved double-bounded pointer comparisons. As mentioned in §6.3, `forward` needs to check whether the object it is considering, which it already knows to be a pointer, is in fact pointing into the `from` space. It uses this function:

```
int Is_from(value * from_start, value * from_limit, value * v) {
  return (from_start <= v && v < from_limit); }
```

Here, the `start` and `limit` pointers are in the same memory block. If `v` is also in the same block, `Is_from` correctly computes whether it is in bounds. However, if `v` is in a different block, the comparison gets stuck rather than returning `false`. Although the Clight code is undefined, we used CompCert's "extcall_properties" to prove that CompCert's compiler transformations will preserve the necessary invariants because the comparison is bounded both above and below (in contrast, single-bounded comparisons need not be semantically preserved in CompCert).

883   The second area of undefined behavior results from our GC's use of the well-established 31-bit
884 integer trick to allow both boxed and unboxed data in block fields [Hickey et al. 2014]. To distinguish
885 them, forward calls Is_block (line 5), which in turn calls

```
int test_int_or_ptr (value x) /* returns 1 if int, 0 if aligned ptr */ {
    return (int)(((intnat)x)&1); }
```

889   When x is an int, this is indeed well-defined, but it is undefined to take the logical and of a pointer,
890 so the code again gets stuck. Here the situation is messier, since CompCert does **not** guarantee
891 that the alignment of a pointer is stable during compilation: in particular, stack-allocated local
892 variables of type char may be packed tightly while assembling the stack frame, thus shifting
893 their alignment. Informally, of course, this cannot occur in the GC since we do not store stack-
894 allocated local variables in the GC-managed heap. We have discussed this issue with the CompCert
895 team [Leroy 2018], and believe that a stronger specification of the extcall_properties (essentially,
896 guaranteeing that non-char pointers do not change alignment) should allow us to prove that
897 CompCert will respect test_int_or_ptr's behavior. The CompCert team understands the
898 issue [?] and wants OCaml-style GCs to have defined behavior in Clight.
899   Other than these two items, the GC is fully defined in Clight—we were even able to prove that
900 all casts (*e.g.* line 6) were well-defined. As a concluding thought, Coq itself is written in OCaml
901 with a similar-style garbage collector. Thus, our GC is at least as well-defined as Coq itself.

7   ENGINEERING OUR TECHNIQUES

904 CompCert is a fully machine-checked verified compiler for C [Leroy 2006]. The Verified Soft-
905 ware Toolchain consists of a series of machine-checked modules written in Coq to reason about
906 (CompCert) C programs [Appel et al. 2014]. Floyd is VST's module for verifying such programs
907 using separation logic. VST's modules interlock so there are no "gaps" in the end-to-end certified
908 results; accordingly all of the rules employed by Floyd have been proved sound with respect to the
909 underlying semantics used by CompCert. Floyd is written in Ltac and Gallina and is designed to
910 help users verify the full functional correctness of their programs. We added two tactics, localize
911 and unlocalize, to integrate the LOCALIZE rule into Floyd as described in §2–§3.

7.1   Localizations in VST with localize and unlocalize

914 Floyd presents users with a pleasant "decorated program" visualization for Hoare proofs, in which
915 users work from the top of the program to the bottom even though the formal proof is maintained
916 as applications of inference rules. For example, suppose the proof goal is $\{P_1\}\ c_1\,;c_2\ \{P_5\}$ and VST's
917 user tells Floyd to apply a Hoare rule for c_1, *e.g.* $\{P_1\}\ c_1\ \{P_2\}$. Floyd will then automatically apply
918 the SEQUENCE rule and show the user $\{P_2\}\ c_2\ \{P_5\}$ as the remaining goal. When the user is in the
919 middle of a verification, the decorated program is partially done (*i.e.* the proof is finished from the
920 top to "the current program point") and the inference tree is also partially done (*i.e.* with holes that
921 are represented by the remaining proof goals in Coq).
922   We wish to preserve this "decorated program" view while extending Floyd to support localization.
923 Our task therefore is to construct a proof in Coq's underlying logic that allows a localization block
924 to be constructed in this manner—that is, we wish to enter a localization block without requiring
925 the user to specify the "exit point" in advance. The engineering is tricky because the proof Floyd
926 is constructing (*i.e.* applications of inference rules) has holes in places where the user's "top to
927 bottom" view of things has not yet arrived.
928   Figure 10 has four partially-decorated "proofs in progress", from both the user's (front end) and
929 Floyd's (back end) points of view. In the first column, from the user's point of view, they saw the
930 assertion P_2 (line 3) and decided to use the localize tactic to zoom into P_3 (line 4). They then

Fig. 10. Front and back ends of `localize` and `unlocalize`

applied some proof rules to move past c_2 to reach the assertion P_4 (line 6). At this point, Floyd does not know when the corresponding `unlocalize` tactic will execute, so it does not know which commands will be inside the block or what the final local and global postconditions will be.

Accordingly, the `localize` tactic builds an incremental proof in the underlying program logic by applying FRAME with an uninstantiated metavariable. The second column of Figure 10 shows the back end with the unknown frame ?F, which will eventually be instantiated by `unlocalize`.

In the third column, the user has advanced past c_3 to reach the local postcondition P_5 and now wishes to `unlocalize` to P_6. Afterwards, the internal state looks like the fourth column, and so to a first approximation, `unlocalize` can instantiate ?F with $P_5 \mathrel{-\!\!*} P_6$. In truth, ?F is chosen more subtly to properly handle both existential variables and modified program variables; `unlocalize` then automatically simplifies the goals to present a cleaner interface to the user. These transformations require the additional theory given in §3.

7.2 Statistics related to our development

All our results in this paper have been machine-checked. Although the size of a development does not perfectly match with that development's importance or implementation difficulty, we present the size nonetheless in Table 1. Our proof script is written in a very dense style. For comparison, verifying a simple 39-line list-based merge sort in VST takes 600 lines. At $\approx$ 400 LOC, the garbage collector is much larger, and is very complicated both mathematically and spatially, in many places teetering on the edge of what can be defined in C. For context, CompCert has 217k LOC, 5,687 definitions, and 6,694 theorems; VST has 623k LOC, 14,038 definitions, and 21,442 theorems.

Component	Section	Files	Size (in lines)	Definitions	Theorems
Common Utilities		10	3,578	44	289
Math Graph Library	§4	20	10,585	216	581
Spatial Graph Library	§5	3	2,328	59	110
Integration into VST	§3,§7	11	2,783	17	172
Marking (graph and DAG)	§3	6	775	9	20
Spanning Tree	§3	5	2,723	17	92
Union-Find (heap and array)	§2	18	3,193	107	135
Garbage Collector	§6	16	13,858	235	712
Total Development		89	39,823	704	2,111

Table 1. Statistics for our code base

8 RELATED WORK

Comparison with [Hobor and Villard 2013]. Our work builds on the theory of ramification by Hobor and Villiard, who verified graph algorithms on pen-and-paper using their RAMIFY rule:

$$\frac{\text{RAMIFY}}{\{L_1\} \ c \ \{L_2\} \qquad G_1 \vdash L_1 * (L_2 \mathrel{-\!\!*} G_2)}{\{G_1\} \ c \ \{G_2\}} \qquad freevars(L_2 \mathrel{-\!\!*} G_2) \cap \mathsf{ModVar}(c) = \emptyset$$

Our LOCALIZE rule upgrades RAMIFY to better handle modified program variables (note the side condition and recall the discussion in §3) and existential quantifiers in postconditions. Hobor and Villard avoided these challenges by proposing a unwieldy variant of RAMIFY called RAMIFYASSIGN, which could reason about the special case of a single assignment x=f(...), assuming the verifier can make the local program translation to x'=f(...); x=x', where x' is fresh. This is nontrivial in large existing formal developments, such as VST, that do not have any way to prove programs equivalent. Hobor and Villard could not verify unmodified program code, modify program variables inside nested localization blocks, or handle multiple assignments in a single block as in lines 14–16 of figure 3. Hobor and Villard avoided existentials in localized postconditions by defining all mathematical operations (*e.g. mark*) as functions rather than as relations; this is fine for pen-and-paper, but painful in a mechanized setting wherein functions must be proven to terminate.

Hobor and Villard treated mathematical graphs as triples (V, E, L) of vertices, edges, and a vertex labeling function; vertices had no more than two neighbors. Our mathematical graph framework (§4) is very modular and general and has been tuned to work smoothly in a mechanized context.

Hobor and Villard erroneously defined spatial graphs recursively (§5.1); unfortunately other members of the research community have since followed them in, *e.g.* [Raad et al. 2015]. We exposed this error and provided a sound and quite general definition for graph (§5.2) that recovers fold/unfold reasoning. We developed a much more general and more modular set of related lemmas and connect our spatial reasoning to the verification framework of CompCert/VST (§7), and our development is entirely machine-checked (§7) whereas they used only pen and paper.

Other verification of graph algorithms and/or ⩊. [Yang 2001]'s verification of the Schorr-Waite algorithm is a landmark in the early separation logic literature. [Krishnaswami 2011] provided the first separation logic proof of union-find. [Bornat et al. 2004] gave an early attempt to reason about graph algorithms in separation logic in a more general way.

[Reynolds 2003] was the first to document the overlapping conjunction ⩊, albeit without any strategy to reason about it using Hoare rules. [Gardner et al. 2012] were the first to reason about a program using ⩊ in Javascript. [Raad et al. 2015] used ⩊ to reason about a concurrent spanning algorithm using a kind of "concurrent localization". [Sergey et al. 2015] also verified a concurrent spanning tree algorithm, and moreover developed mechanized Coq proofs.

A decade after Yang verified Schorr-Waite on paper, [Leino 2010] automated its verification.

Verification tools. Our work interacts with the Floyd [Appel et al. 2014] verification tool. Charge! likewise uses Coq tactics to work with a shallow embedding of higher order separation logic, but focuses on OO programs written in Java/C# [Bengtson et al. 2012]. Iris Proof Mode provides a similar framework for higher-order concurrent reasoning [Krebbers et al. 2017]. A more automated approach to verification of low level programs using Coq is the Bedrock framework [Chlipala 2011].

Many automated verification tools also use separation logic in a forward reasoning style as does HIP/SLEEK, including Smallfoot [Berdine et al. 2005], jStar [Distefano and Parkinson 2008], and Verifast [Jacobs et al. 2011]. One of HIP/SLEEK's distinguishing features is good support for user-defined inductive predicates rather than a library of pre-defined predicates for lists, trees etc.

Dafny [Leino 2010] and KeY [Beckert et al. 2007] are verifiers not based on separation logic; KeY uses an interactive verifier while Dafny pursues automation with Z3 [de Moura and Bjørner 2008].

Mechanized mathematical graph theory. There is a long history, going back at least 25 years, of mechanized reasoning about mathematical graphs [Wong 1991]. The most famous mechanically verified "graph theorem" is the Four Color Theorem [Gonthier 2005]; however the development actually uses hypermaps instead of graphs. Noschinski built a graph library in Isabelle/HOL whose formalization is the closest to ours [Noschinski 2015b], *e.g.* supporting graphs with labeled and parallel arcs. [Dubois et al. 2015; Noschinski 2015a] used proof assistants to design verifiable checkers for solutions to graph problems. [Bauer and Nipkow 2002; Yamamoto et al. 1995] use an inductive encoding of graphs to formalize planar graph theory.

[Guéneau et al. 2018] formalised "time credits" in Separation Logic and big-*O* notation in Coq and then verified both the correctness and the time complexity of the union-find data structure.

Verification of garbage collection algorithms. Schism [Gammie et al. 2015; Pizlo et al. 2010] is a certified concurrent collector built in a Java VM that services multi-core architectures with weak memory consistency. McCreight *et al.* [McCreight et al. 2010, 2007] introduce GCminor, which is a certified translation step added to CompCert's translation from Clight to assembly. GCminor makes explicit the specific invariants that the garbage collector relies upon, thus minimising errors due to the violation of invariants between the garbage collector and the mutator. Hawblitzel and Petrank [Petrank and Hawblitzel 2010] annotate x86 code for two GCs by hand, and then use Boogie and the Z3 automated theorem prover to verify their correctness automatically.

The closest piece of work to our certified GC is probably the excellent certified GC for the Cake ML project [Ericsson et al. 2017], since both integrate a certified GC into a certified compiler for a functional language. Their GC is written closer to assembly than C, which is both a positive—in that they avoid undefined behaviors—and a negative, in that their GC is harder to understand and upgrade and cannot take advantage of the mature CompCert compiler. Their GC lacks some of our optimisations (*e.g.* they have only three generations), but on the other hand handles mutation in the GC heap. The largest difference, however, is that we present an integrated graph framework suitable for reasoning about many graph algorithms, of which our GC is merely the flagship. In contrast, they focus much more narrowly on the problem of certified GCs.

9 FUTURE WORK AND CONCLUSION

In the future we plan to improve the pure reasoning of graphs and similar data structures, in particular to add automation. We have also begun to investigate integrating our techniques into the HIP/SLEEK toolchain [Chin et al. 2010], which as compared to VST provides more automation at the cost of lower expressivity. We are also interested in investigating better ways to handle the kinds of undefined behavior upon which real C systems code sometimes relies.

Our main contributions were as follows. We developed a mathematical graph library that was powerful enough to reason about graph-manipulating algorithms written in real C code. We connected these mathematical graphs to spatial graphs in the heap via separation logic. We developed localization blocks to smoothly reason about a local action's effect on a global context in a mechanized context, including a robust treatment of modified program variables and existential quantifiers in postconditions. We demonstrated our techniques on several nontrivial examples, including union-find and spanning tree. Our flagship example is the verification of the garbage collector for the CertiCoq project, during which we found two places in which the C semantics is too weak to define an OCaml-style GC. We integrated our techniques into the VST toolset.

REFERENCES

Abhishek Anand, Andrew Appel, Greg Morrisett, Zoe Paraskevopoulou, Randy Pollack, Olivier Savary Belanger, Matthieu Sozeau, and Matthew Weaver. 2017. CertiCoq: A verified compiler for Coq. In *The Third International Workshop on Coq for Programming Languages (CoqPL)*.

Andrew W. Appel, Robert Dockins, Aquinas Hobor, Lennart Beringer, Josiah Dodds, Gordon Stewart, Sandrine Blazy, and Xavier Leroy. 2014. *Program Logics for Certified Compilers*. Cambridge University Press, New York, NY, USA.

Andrew W. Appel and David McAllester. 2001. An Indexed Model of Recursive Types for Foundational Proof-Carrying Code. *ACM Transactions on Programming Languages and Systems* 23(5) (2001), 657–683.

Andrew W. Appel, Paul-André Melliès, Christopher D. Richards, and Jerôme Vouillon. 2007. A Very Modal Model of a Modern, Major, General Type System. In *Proceedings of the 34th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'07)*. 109–122.

Gertrud Bauer and Tobias Nipkow. 2002. The 5 colour theorem in Isabelle/Isar. In *International Conference on Theorem Proving in Higher Order Logics*. Springer, 67–82.

Bernhard Beckert, Reiner Hähnle, and Peter H. Schmitt. 2007. *Verification of Object-oriented Software: The KeY Approach*. Springer-Verlag, Berlin, Heidelberg.

Jesper Bengtson, Jonas Braband Jensen, and Lars Birkedal. 2012. Charge! - A Framework for Higher-Order Separation Logic in Coq. In *Interactive Theorem Proving - Third International Conference, ITP 2012, Princeton, NJ, USA, August 13-15, 2012. Proceedings*. 315–331.

Josh Berdine, Cristiano Calcagno, and Peter W. O'Hearn. 2005. Smallfoot: Modular Automatic Assertion Checking with Separation Logic. In *FMCO*. 115–137.

R. Bornat, C. Calcagno, and P. O'Hearn. 2004. Local reasoning, separation and aliasing. In *SPACE*, Vol. 4.

Richard Bornat, Cristiano Calcagno, and Hongseok Yang. 2006. Variables as Resource in Separation Logic. *ENTCS* 155 (2006), 247–276.

C. J. Cheney. 1970. A nonrecursive list compacting algorithm. *Commun. ACM* 13, 11 (1970), 677–678. https://doi.org/10.1145/362790.362798

Wei Ngan Chin, Cristina David, Huu Hai Nguyen, and Shengchao Qin. 2010. Automated verification of shape, size and bag properties via user-defined predicates in separation logic. *Science of Computer Programming* 77(9) (2010), 1,006–1,036.

Adam Chlipala. 2011. Mostly-automated verification of low-level programs in computational separation logic. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2011, San Jose, CA, USA, June 4-8, 2011*. 234–245.

Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford S. Stein. 2009. *Introduction to algorithms, 3rd edition*. MIT Press and McGraw-Hill.

Leonardo Mendonça de Moura and Nikolaj Bjørner. 2008. Z3: An Efficient SMT Solver. In *Tools and Algorithms for the Construction and Analysis of Systems, 14th International Conference, TACAS 2008, Held as Part of the Joint European Conferences on Theory and Practice of Software, ETAPS 2008, Budapest, Hungary, March 29-April 6, 2008. Proceedings*. 337–340. https://doi.org/10.1007/978-3-540-78800-3_24

Dino Distefano and Matthew J. Parkinson. 2008. jStar: towards practical verification for java. In *Proceedings of the 23rd Annual ACM SIGPLAN Conference on Object-Oriented Programming, Systems, Languages, and Applications, OOPSLA 2008, October 19-23, 2008, Nashville, TN, USA*. 213–226. https://doi.org/10.1145/1449764.1449782

Catherine Dubois, Sourour Elloumi, Benoit Robillard, and Clément Vincent. 2015. Graphes et couplages en Coq. In *Vingt-sixièmes Journées Francophones des Langages Applicatifs (JFLA 2015)*.

Adam Sandberg Ericsson, Magnus O. Myreen, and Johannes Åman Pohjola. 2017. A Verified Generational Garbage Collector for CakeML. In *Interactive Theorem Proving - 8th International Conference, ITP 2017, Brasília, Brazil, September 26-29, 2017, Proceedings*. 444–461. https://doi.org/10.1007/978-3-319-66107-0_28

Peter Gammie, Antony L Hosking, and Kai Engelhardt. 2015. Relaxing safely: verified on-the-fly garbage collection for x86-TSO. In *ACM SIGPLAN Notices*, Vol. 50. ACM, 99–109.

Philippa Gardner, Sergio Maffeis, and Gareth David Smith. 2012. Towards a program logic for JavaScript. In *Proceedings of the 39th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2012, Philadelphia, Pennsylvania, USA, January 22-28, 2012*. 31–44. https://doi.org/10.1145/2103656.2103663

Georges Gonthier. 2005. A computer-checked proof of the four colour theorem.

Armaël Guéneau, Arthur Charguéraud, and François Pottier. 2018. A Fistful of Dollars: Formalizing Asymptotic Complexity Claims via Deductive Program Verification. In *Programming Languages and Systems - 27th European Symposium on Programming, ESOP 2018, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2018, Thessaloniki, Greece, April 14-20, 2018, Proceedings*. 533–560. https://doi.org/10.1007/978-3-319-89884-1_19

Jason Hickey, Anil Madhavapeddy, and Yaron Minsky. 2014. *Real World OCaml*. OReilly.

Aquinas Hobor and Jules Villard. 2013. The ramifications of sharing in data structures. In *Proceedings of the 40th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL'13)*. 523–536.

Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and Frank Piessens. 2011. VeriFast: A Powerful, Sound, Predictable, Fast Verifier for C and Java. In *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA, April 18-20, 2011. Proceedings*. 41–55.

johndoe20190406. 2019. johndoe20190406/RamifyCoq. https://github.com/johndoe20190406/RamifyCoq

Robbert Krebbers, Amin Timany, and Lars Birkedal. 2017. Interactive proofs in higher-order concurrent separation logic. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL 2017, Paris, France, January 18-20, 2017*. 205–217. http://dl.acm.org/citation.cfm?id=3009855

Neelakantan R. Krishnaswami. 2011. *Verifying Higher-Order Imperative Programs with Higher-Order Separation Logic*. Ph.D. Dissertation.

K. Rustan M. Leino. 2010. Dafny: An Automatic Program Verifier for Functional Correctness. In *Logic for Programming, Artificial Intelligence, and Reasoning - 16th International Conference, LPAR-16, Dakar, Senegal, April 25-May 1, 2010, Revised Selected Papers*. 348–370. https://doi.org/10.1007/978-3-642-17511-4_20

Xavier Leroy. 2006. Formal certification of a compiler back-end or: programming a compiler with a proof assistant. In *Proceedings of the 33rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2006, Charleston, South Carolina, USA, January 11-13, 2006*. 42–54.

Xavier Leroy. 2018. Re: The dark corners of C light pointer semantics.

Andrew McCreight, Tim Chevalier, and Andrew Tolmach. 2010. A certified framework for compiling and executing garbage-collected languages. In *ACM Sigplan Notices*, Vol. 45. ACM, 273–284.

Andrew McCreight, Zhong Shao, Chunxiao Lin, and Long Li. 2007. A general framework for certifying garbage collectors and their mutators. In *ACM SIGPLAN Notices*, Vol. 42. ACM, 468–479.

Lars Noschinski. 2015a. *Formalizing Graph Theory and Planarity Certificates*. Ph.D. Dissertation. Universität München.

Lars Noschinski. 2015b. A Graph Library for Isabelle. *Mathematics in Computer Science* 9, 1 (2015), 23–39. https://doi.org/10.1007/s11786-014-0183-z

Erez Petrank and Chris Hawblitzel. 2010. Automated Verification of Practical Garbage Collectors. *Logical Methods in Computer Science* 6 (2010).

Filip Pizlo, Lukasz Ziarek, Petr Maj, Antony L Hosking, Ethan Blanton, and Jan Vitek. 2010. Schism: fragmentation-tolerant real-time garbage collection. *ACM Sigplan Notices* 45, 6 (2010), 146–159.

Azalea Raad, Jules Villard, and Philippa Gardner. 2015. CoLoSL: Concurrent Local Subjective Logic. In *Programming Languages and Systems - 24th European Symposium on Programming, ESOP 2015, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2015, London, UK, April 11-18, 2015. Proceedings*. 710–735. https://doi.org/10.1007/978-3-662-46669-8_29

John C. Reynolds. 2003. A Short Course on Separation Logic. (2003). http://www.cs.cmu.edu/afs/cs.cmu.edu/project/fox-19/member/jcr/wwwaac2003/notes7.ps.

Ilya Sergey, Aleksandar Nanevski, and Anindya Banerjee. 2015. Mechanized Verification of Fine-grained Concurrent Programs. In *PLDI*. 77–87.

Alfred Tarski. 1955. A lattice-theoretical fixpoint theorem and its applications. *Pacific Journal of Mathematics* 5 (1955), 285–309.

Wai Wong. 1991. A Simple Graph Theory And Its Application In Railway Signaling. In *HOL Theorem Proving System and Its Applications, 1991., International Workshop on the*. 395–409. https://doi.org/10.1109/HOL.1991.596304

Mitsuharu Yamamoto, Shin-ya Nishizaki, Masami Hagiya, and Yozo Toda. 1995. Formalization of planar graphs. In *International Conference on Theorem Proving in Higher Order Logics*. Springer, 369–384.

Hongseok Yang. 2001. *Local Reasoning for Stateful Programs*. Ph.D. Dissertation. University of Illinois.

A SPANNING AND COPYING

In Figure 11 we show a simplified proof script for the spanning tree algorithm. Unlike graph marking, the spanning tree algorithm changes the structure of the graph, leading to a more complicated specification, in both the pure part and the spatial part. Observe that the *span* relation is rather long; the *e_span* handles the case of either calling spanning tree or deleting an edge.

We put the proof sketch of the graph copying algorithm in Figure 12 and Figure 13. Just like other parts of the paper, both algorithms have been machine verified.

B DIFFICULTY USING graph$_T$

See Figure 14 for an attempt to prove the entailment $100 \mapsto 42, 100, 0 \vdash \text{graph}_T(100, \hat{\gamma})$. Part of the problem is that the recursive structure interacts very badly with $\uplus$: if the recursion involved $*$ then it **would** be provable, by induction on the finite memory (each "recursive call" would be on a strictly smaller subheap). This is why Knaster-Tarski works so well with list, tree, and DAG predicates in separation logic.

C PROBLEM WITH APPEL AND MCALLESTER'S FIXPOINT

Appel and McAllester proposed another fixpoint μ_A that is sometimes used to define recursive predicates in separation logic [Appel and McAllester 2001]. This time the functional F_P needs to be *contractive*, which to a first order of approximation means that all recursion needs to be guarded by the "approximation modality" $\triangleright$ [Appel et al. 2007], *i.e.* our graph predicate would look like

$$\text{graph}_A(x, \gamma) \overset{\Delta}{=}$$
$$(x = 0 \wedge \text{emp}) \vee \exists m, l, r.\ \gamma(x) = (m, l, r) \wedge$$
$$x \mapsto m, l, r \uplus \triangleright\text{graph}_A(l, \gamma) \uplus \triangleright\text{graph}_A(r, \gamma)$$

Unfortunately, $\triangleright P$ is not precise for all P, so graph_A is not precise either. The approximation modality's universal imprecision has never been noticed before.

D STRUCTURE OF THE GARBAGE COLLECTOR PROGRAM

CertiCoq uses a generational copying garbage collector that is inspired by the OCaml GC. The heap is divided into a series of disjoint spaces called *generations*. The size of the first generation is carefully calculated, and then subsequent generations double in size. The mutator only ever allocates new memory in the first, smallest generation of heap, which is called the nursery. If it finds that the nursery is full, the mutator calls the GC to free up space. The GC collects the nursery (now called the *from* generation) into the second generation (the *to* generation): it examines the elements in *from*, sees if they are accessible by the mutator, and, if they are, copies them over to *to*. This copying is achieved over a few steps, and we will explain these shortly, but the larger picture is that everything of importance in *from* gets copied to *to*, and so *from* can safely be reset.

An important subtlety here is that *to* had enough room to accept *from*'s items. In the (empirically improbable) worst case, *all* of *from*'s fields were copied over to *to*. Because *to* has twice the capacity of *from*, *to* could not have been more than half full when the collection started. This guarantee must be renewed before the next collection. So, in case the collection of the nursery caused the second generation to become more than half full, the second generation is collected into the third. This makes both the first and second generations empty, thus ensuring the guarantee trivially. It should be clear to see that this may also trigger further collections in a cascade effect. The GC's task is only complete once this cascade (if any) is over. It returns control to the mutator, which goes ahead with the allocation that it was trying to perform in the nursery.

Having shown that the overall collection works via (a series of) two-generational collections, we now zoom in and explain a two-generation collection. The GC starts at the mutator-owned arguments array, whose fields are either data, or pointers that point at memory blocks in the heap. It ignores the data entirely, and, among the pointers, cares only for the pointers that point into the *from* generation. For each pointer that points into *from*, it copies its target block to *to*, simply adding it in its entirety after *to*'s last-used memory field, which is called *next*.

This operation only takes care of the blocks in the heap that the arguments array was pointing at directly, so the GC still has to copy over indirectly-accessible blocks. Of course, the only way to access an indirect block is via one of the direct blocks that it has just finished copying into a

contiguous array. It starts at the old *next* in *to* and works its way "upwards" through the freshly copied blocks, again looking exclusively for pointers that point into *from* and copying over their target blocks into *to*. In the *to* generation, these newly copied blocks simply get stacked atop our first batch of copied blocks.

The mutator's dependency graph has indefinite depth, so the second batch of copied blocks may still have pointers into *from*. However, thanks to this systematic copying strategy, it is very easy to take care of all indirect blocks. The GC simply keeps scanning upwards in *to*, copying over blocks from *from* as necessary, until the scanning pointer catches up to the last-used field in *to*. This completes a collection from *from* to *to*, copying all blocks that lived in *from* and were of interest to the mutator. *from* is now reset.

A good question at this juncture is why this rather selective scan of the args array and the heap is good enough to collect *from*. The GC definitely collected every direct block by scanning the args array, but what of the indirect blocks? Couldn't there be valid indirect links that start either below *next* in the *to* generation, or from other generations altogether?

Both of these turn out to be impossible because the mutator behaves in a purely functional manner. The heap is chronologically faithful, in that higher-indexed generations host objects that were allocated earlier. Because of the immutability of objects in a purely functional language, it is impossible for objects to point "backwards" to a younger generation, as the older object would not have had known about the younger at the time of its allocation, and could not have been modified after its allocation. Fields living in generations younger than *from* can point into *from* during normal mutator activity, but this is impossible at the time of collection: *from* is only ever collected either if it is the nursery or if a cascade effect has caused all generations younger than it to be collected and reset. In fact, the only time the GC ever sees backwards pointers is when it creates (and quickly fixes) them during during its activities.

E   CODE FOR FORWARD RELATION

```
1   Inductive
2 forward_relation (from to : nat) : nat -> forward_t -> LGraph -> LGraph -> Prop :=
3     fr_z : forall (depth : nat) (z : Z) (g : LGraph),
4           forward_relation from to depth (inl (inl (inl z))) g g
5   | fr_p : forall (depth : nat) (p : GC_Pointer) (g : LGraph),
6           forward_relation from to depth (inl (inl (inr p))) g g
7   | fr_v_not_in : forall (depth : nat) (v : VType) (g : LGraph),
8                 vgeneration v <> from ->
9                 forward_relation from to depth (inl (inr v)) g g
10  | fr_v_in_forwarded : forall (depth : nat) (v : VType)
11                          (g : LabeledGraph VType EType raw_vertex_block unit
12                               graph_info),
13                  vgeneration v = from ->
14                  raw_mark (vlabel g v) = true ->
15                  forward_relation from to depth (inl (inr v)) g g
16  | fr_v_in_not_forwarded_O : forall (v : VType)
17                          (g : LabeledGraph VType EType raw_vertex_block unit
18                                graph_info),
19                  vgeneration v = from ->
20                  raw_mark (vlabel g v) = false ->
21                  forward_relation from to 0
22                    (inl (inr v)) g (lgraph_copy_v g v to)
23  | fr_v_in_not_forwarded_Sn : forall (depth : nat) (v : VType)
24                          (g : LabeledGraph VType EType raw_vertex_block unit
25                                graph_info) (g' : LGraph),
26                  vgeneration v = from ->
27                  raw_mark (vlabel g v) = false ->
28                  let new_g := lgraph_copy_v g v to in
29                  forward_loop from to depth
30                    (vertex_pos_pairs new_g (new_copied_v g to)) new_g
31                    g' ->
32                  forward_relation from to (S depth) (inl (inr v)) g g'
```

```
33    | fr_e_not_to : forall (depth : nat) (e : EType) (g : LGraph),
34                   vgeneration (dst (pg_lg g) e) <> from ->
35                   forward_relation from to depth (inr e) g g
36    | fr_e_to_forwarded : forall (depth : nat) (e : EType) (g : LGraph),
37                   vgeneration (dst (pg_lg g) e) = from ->
38                   raw_mark (vlabel g (dst (pg_lg g) e)) = true ->
39                   let new_g :=
40                     labeledgraph_gen_dst g e
41                       (copied_vertex (vlabel g (dst (pg_lg g) e))) in
42                   forward_relation from to depth (inr e) g new_g
43    | fr_e_to_not_forwarded_O : forall (e : EType) (g : LGraph),
44                   vgeneration (dst (pg_lg g) e) = from ->
45                   raw_mark (vlabel g (dst (pg_lg g) e)) = false ->
46                   let new_g :=
47                     labeledgraph_gen_dst
48                       (lgraph_copy_v g (dst (pg_lg g) e) to) e
49                       (new_copied_v g to) in
50                   forward_relation from to 0 (inr e) g new_g
51    | fr_e_to_not_forwarded_Sn : forall (depth : nat) (e : EType) (g g' : LGraph),
52                   vgeneration (dst (pg_lg g) e) = from ->
53                   raw_mark (vlabel g (dst (pg_lg g) e)) = false ->
54                   let new_g :=
55                     labeledgraph_gen_dst
56                       (lgraph_copy_v g (dst (pg_lg g) e) to) e
57                       (new_copied_v g to) in
58                   forward_loop from to depth
59                     (vertex_pos_pairs new_g (new_copied_v g to)) new_g
60                     g' ->
61                   forward_relation from to (S depth) (inr e) g g'
62    with forward_loop (from to : nat)
63         : nat -> list forward_p_type -> LGraph -> LGraph -> Prop :=
64    fl_nil : forall (depth : nat) (g : LGraph), forward_loop from to depth [] g g
65    | fl_cons : forall (depth : nat) (g1 g2 g3 : LGraph) (f : forward_p_type)
66                   (fl : list forward_p_type),
67                 forward_relation from to depth (forward_p2forward_t f [] g1) g1 g2 ->
68                 forward_loop from to depth fl g2 g3 ->
69                 forward_loop from to depth (f :: fl) g1 g3
```

```c
1  struct Node {
2      int m;
3      struct Node * l;
4      struct Node * r; };
```
5 `// We use` R `to represent` reachable(γ, x)
```c
6  void spanning(struct Node * x) {
```
7 `//` {graph(x, γ) $\land$ γ(x).1 = 0}
```c
8      struct Node * l, * r; int root_mark;
```
9 `//` {graph(x, γ) $\land$ $\exists l, r.\ \gamma$(x) = (0, l, r)}
10 `//` {graph(x, γ) $\land$ γ(x) = (0, l, r)}
11 `//` {vertices_at(reachable(γ, x), γ) $\land$ γ(x) = (0, l, r)}
12 `//` {vertices_at(R, γ) $\land$ γ(x) = (0, l, r)}
13 `//` $\searrow$ {x $\mapsto$ 0, l, r $\land$ γ(x) = (0, l, r)}
```c
14     l = x -> l; r = x -> r; x -> m = 1;
```
15 `//` $\nearrow$ {x $\mapsto$ 1, l, r $\land$ γ(x) = (0, l, r) $\land$ $\exists\gamma_1.\ $ mark1(γ, x, γ_1)}
16 `//` {$\exists\gamma_1.\ $ vertices_at(R, γ_1) $\land$ γ(x) = (0, l, r) $\land$ mark1(γ, x, γ_1)}
17 `//` {vertices_at(R, γ_1) $\land$ γ(x) = (0, l, r) $\land$ mark1(γ, x, γ_1)}
```c
18     if (l) {
19         root_mark = l -> m;
20         if (root_mark == 0) {
21             spanning(l);
22         } else { x -> l = 0; } }
```
23 `//` $\left\{ \begin{array}{l} \exists\gamma_2.\ \text{vertices\_at}(R, \gamma_2) \land \gamma(\text{x}) = (0, \text{l}, \text{r}) \land \\ \textit{mark1}(\gamma, \text{x}, \gamma_1) \land \textit{e\_span}(\gamma_1, \text{x.L}, \gamma_2) \end{array} \right\}$
24 `//` $\left\{ \begin{array}{l} \text{vertices\_at}(R, \gamma_2) \land \gamma(\text{x}) = (0, \text{l}, \text{r}) \land \\ \textit{mark1}(\gamma, \text{x}, \gamma_1) \land \textit{e\_span}(\gamma_1, \text{x.L}, \gamma_2) \end{array} \right\}$
```c
25     if (r) {
26         root_mark = r -> m;
27         if (root_mark == 0) {
28             spanning(r);
29         } else { x -> r = 0; } }
```
30 `//` $\left\{ \begin{array}{l} \exists\gamma_3.\ \text{vertices\_at}(R, \gamma_3) \land \gamma(\text{x}) = (0, \text{l}, \text{r}) \land \\ \textit{mark1}(\gamma, \text{x}, \gamma_1) \land \textit{e\_span}(\gamma_1, \text{x.L}, \gamma_2) \land \textit{e\_span}(\gamma_2, \text{x.R}, \gamma_3) \end{array} \right\}$
31 `}` `//` {$\exists\gamma_3.\ $ vertex_at(reachable(γ, x), γ_3) $\land$ span(γ, x, γ_3)}

$$\text{vertices\_at(reachable}(\gamma_1, \text{x}), \gamma_2) \stackrel{\Delta}{=} \underset{v\in\text{reachable}(\gamma_1,\text{x})}{\bigstar} v \mapsto \gamma_2(v)$$

$$span(\gamma_1, \text{x}, \gamma_2) \stackrel{\Delta}{=} mark(\gamma_1, \text{x}, \gamma_2) \land \gamma_1 \uparrow (\lambda v.\text{x} \overset{\gamma_1}{\leadsto_0^\star} v) \text{ is a tree} \land$$

$$\gamma_1 \uparrow (\lambda v.\neg\text{x} \overset{\gamma_1}{\leadsto_0^\star} v) = \gamma_2 \uparrow (\lambda v.\neg\text{x} \overset{\gamma_1}{\leadsto_0^\star} v) \land$$

$$(\forall v.\ \text{x} \overset{\gamma_1}{\leadsto_0^\star} v \Rightarrow \gamma_2 \models \text{x} \leadsto v) \land$$

$$(\forall a, b.\ \text{x} \overset{\gamma_1}{\leadsto_0^\star} a \Rightarrow \neg\text{x} \overset{\gamma_1}{\leadsto_0^\star} b \Rightarrow \neg\gamma_2 \models a \leadsto b)$$

$$e\_span(\gamma_1, e, \gamma_2) \stackrel{\Delta}{=} \begin{cases} \gamma_1 - e = \gamma_2 & t(\gamma_1, e) = 1 \\ span(\gamma_1, t(\gamma_1, e), \gamma_2) & t(\gamma_1, e) = 0 \end{cases}$$

Fig. 11. Clight code and proof sketch for bigraph spanning tree.

```
1  struct Node {
2      int m;
3      struct Node * l;
4      struct Node * r;  };
```
5 `// We use` $x \overset{\gamma}{\longleftrightarrow} x'$ `to represent` $x = x' = 0 \vee \gamma(x) = (x', \_, \_)$
```
6  struct Node * copy(struct Node * x) {
7      struct Node * l, * r, * x0, * l0, * r0;
```
8 `//` $\{\text{graph}(x, \gamma)\}$
```
9      if (x == 0)
10         return 0;
```
11 `//` $\{\text{graph}(x, \gamma) \wedge x \neq 0\}$

12 `//` $\{\text{graph}(x, \gamma) \wedge \exists x_0, l, r. \ \gamma(x) = (x_0, l, r)\}$

13 `//` $\{\text{graph}(x, \gamma) \wedge \gamma(x) = (x_0, l, r)\}$

14 `//` $\searrow \{x \mapsto x_0, l, r \wedge \gamma(x) = (x_0, l, r)\}$
```
15     x0 = x -> m;
```
16 `//` $\nearrow \{x \mapsto x_0, l, r \wedge \gamma(x) = (x_0, l, r) \wedge x0 = x_0\}$

17 `//` $\{\text{graph}(x, \gamma) \wedge \gamma(x) = (x0, l, r)\}$
```
18     if (x0 != 0)
19         return x0;
```
20 `//` $\{\text{graph}(x, \gamma) \wedge \gamma(x) = (0, l, r)\}$
```
21     x0 = (struct Node *) mallocN (sizeof (struct Node));
```
22 `//` $\{\text{graph}(x, \gamma) * x0 \mapsto \_, \_, \_ \wedge \gamma(x) = (0, l, r)\}$

23 `//` $\searrow \{x \mapsto 0, l, r * x0 \mapsto \_, \_, \_ \wedge \gamma(x) = (0, l, r)\}$
```
24     l = x -> l; r = x -> r; x -> m = x0; x0 -> m = 0;
```
25 `//` $\nearrow \left\{ \begin{array}{l} x \mapsto x0, l, r * x0 \mapsto 0, \_, \_ \wedge \\ \gamma(x) = (0, l, r) \wedge \exists \gamma_1 \gamma_1'. v\_copy1(\gamma, x, \gamma_1, \gamma_1') \end{array} \right\}$

26 `//` $\left\{ \begin{array}{l} \exists \gamma_1 \gamma_1'. \text{graph}(x, \gamma_1) * x0 \mapsto 0, \_, \_ \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \end{array} \right\}$

27 `//` $\left\{ \begin{array}{l} \text{graph}(x, \gamma_1) * x0 \mapsto 0, \_, \_ \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \end{array} \right\}$

28 `//` $\left\{ \begin{array}{l} \text{graph}(x, \gamma_1) * x0 \mapsto 0, \_, \_ * \text{holegraph}(x0, \gamma_1') \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \end{array} \right\}$

29 `//` $\searrow \{\text{graph}(l, \gamma_1)\}$
```
30     l0 = copy(l);
```
31 `//` $\nearrow \left\{ \begin{array}{l} \exists \gamma_2 \gamma_2''. \text{graph}(l, \gamma_2) * \text{graph}(l0, \gamma_2'') \wedge \\ copy(\gamma_1, l, \gamma_2, \gamma_2'') \wedge l \overset{\gamma_2}{\longleftrightarrow} l0 \end{array} \right\}$

32 `//` $\left\{ \begin{array}{l} \exists \gamma_2 \gamma_2''. \text{graph}(x, \gamma_2) * x0 \mapsto 0, \_, \_ * \text{holegraph}(x0, \gamma_1') * \\ \text{graph}(l0, \gamma_2'') \wedge \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \wedge \\ copy(\gamma_1, l, \gamma_2, \gamma_2'') \wedge l \overset{\gamma_2}{\longleftrightarrow} l0 \end{array} \right\}$

33 `//` $\left\{ \begin{array}{l} \exists \gamma_2 \gamma_2'. \text{graph}(x, \gamma_2) * x0 \mapsto 0, \_, \_ * \text{holegraph}(x0, \gamma_2') \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \wedge \\ e\_copy(\gamma_1, \gamma_1', x.L, \gamma_2, \gamma_2') \wedge l \overset{\gamma_2}{\longleftrightarrow} l0 \end{array} \right\}$

34 `//` $\left\{ \begin{array}{l} \text{graph}(x, \gamma_2) * x0 \mapsto 0, \_, \_ * \text{holegraph}(x0, \gamma_2') \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \wedge \\ e\_copy(\gamma_1, \gamma_1', x.L, \gamma_2, \gamma_2') \wedge l \overset{\gamma_2}{\longleftrightarrow} l0 \end{array} \right\}$
```
35     x0 -> l = l0;
```
36 `//` $\left\{ \begin{array}{l} \text{graph}(x, \gamma_2) * x0 \mapsto 0, l0, \_ * \text{holegraph}(x0, \gamma_2') \wedge \\ \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \wedge \\ e\_copy(\gamma_1, \gamma_1', x.L, \gamma_2, \gamma_2') \wedge l \overset{\gamma_2}{\longleftrightarrow} l0 \end{array} \right\}$

37 `//` $\searrow \{\text{graph}(r, \gamma_2)\}$
```
38     r0 = copy(r);
```
39 `//` $\nearrow \left\{ \begin{array}{l} \exists \gamma_3 \gamma_3''. \text{graph}(r, \gamma_3) * \text{graph}(r0, \gamma_3'') \wedge \\ copy(\gamma_2, r, \gamma_3, \gamma_3'') \wedge r \overset{\gamma_3}{\longleftrightarrow} r0 \end{array} \right\}$

40 `//` $\left\{ \begin{array}{l} \exists \gamma_3 \gamma_3''. \text{graph}(x, \gamma_3) * x0 \mapsto 0, l0, \_ * \text{holegraph}(x0, \gamma_3') * \\ \text{graph}(r0, \gamma_3'') \wedge \gamma(x) = (0, l, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma_1') \wedge \\ e\_copy(\gamma_1, \gamma_1', x.L, \gamma_2, \gamma_2') \wedge copy(\gamma_2, r, \gamma_3, \gamma_3'') \wedge \\ l \overset{\gamma_2}{\longleftrightarrow} l0 \wedge r \overset{\gamma_3}{\longleftrightarrow} r0 \end{array} \right.$

1 // $\left\{ \begin{array}{l} \text{graph}(x, \gamma_3) * x0 \mapsto 0, \, l0, \, \_ * \text{holegraph}(x0, \gamma'_3) \wedge \\ \gamma(x) = (0, \, l, \, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma'_1) \wedge \\ e\_copy(\gamma_1, \gamma'_1, x.L, \gamma_2, \gamma'_2) \wedge e\_copy(\gamma_2, \gamma'_2, x.R, \gamma_3, \gamma'_3) \wedge \\ l \overset{\gamma_2}{\longleftrightarrow} l0 \wedge r \overset{\gamma_3}{\longleftrightarrow} r0 \end{array} \right\}$

2   $x0 \; \mathtt{->} \; r \; = \; r0;$

3 // $\left\{ \begin{array}{l} \text{graph}(x, \gamma_3) * x0 \mapsto 0, \, l0, \, r0 * \text{holegraph}(x0, \gamma'_3) \wedge \\ \gamma(x) = (0, \, l, \, r) \wedge v\_copy1(\gamma, x, \gamma_1, \gamma'_1) \wedge \\ e\_copy(\gamma_1, \gamma'_1, x.L, \gamma_2, \gamma'_2) \wedge e\_copy(\gamma_2, \gamma'_2, x.R, \gamma_3, \gamma'_3) \wedge \\ l \overset{\gamma_2}{\longleftrightarrow} l0 \wedge r \overset{\gamma_3}{\longleftrightarrow} r0 \end{array} \right\}$

4 // $\left\{ \text{graph}(x, \gamma_3) * \text{graph}(x0, \gamma'_3) \wedge copy(\gamma, x, \gamma_3, \gamma'_3) \wedge x \overset{\gamma_3}{\longleftrightarrow} x0 \right\}$

$$\text{holegraph}(x, \gamma) \overset{\Delta}{=} \underset{v \in \text{reachable}(\gamma, x) - \{x\}}{\bigstar} v \mapsto \gamma(v)$$

$$iso(f_V, f_E, \gamma_1, \gamma_2) \overset{\Delta}{=} f_V \text{ is a bijection between } \phi_V(\gamma_1) \text{ and } \phi_V(\gamma_2) \wedge$$
$$f_E \text{ is a bijection between } \phi_E(\gamma_1) \text{ and } \phi_E(\gamma_2) \wedge$$
$$\forall e, f_V(s(\gamma_1, e)) = s(\gamma_2, f_E(e)) \wedge$$
$$\forall e, f_V(d(\gamma_1, e)) = d(\gamma_2, f_E(e))$$

$$v\_copy1(\gamma_1, x, \gamma_2, \gamma'_2) \overset{\Delta}{=} \exists x'. x \neq 0 \wedge mark1(\gamma_1, x, \gamma_2) \wedge$$
$$x \overset{\gamma_2}{\longleftrightarrow} x' \wedge \gamma'_2 = \{x_0\}$$

$$copy(\gamma_1, x, \gamma_2, \gamma'_2) \overset{\Delta}{=} mark(\gamma_1, x, \gamma_2) \wedge$$
$$\exists f_V f_E. iso(f_V, f_E, \gamma_2 \uparrow (\lambda v. x \overset{\gamma_2}{\leadsto_0^\star} v), \gamma'_2) \wedge$$
$$\forall x x'. f_V(x) = x' \Leftrightarrow x \overset{\gamma_2}{\longleftrightarrow} x'$$

$$e\_copy(\gamma_1, \gamma'_1, e, \gamma_2, \gamma'_2) \overset{\Delta}{=} \exists \gamma''_2. \gamma'_2 = \gamma'_1 + \gamma''_2 \wedge$$
$$mark(\gamma_1, x, \gamma_2) \wedge \exists f_V f_E.$$
$$iso(f_V, f_E, \{e\} + \gamma_2 \uparrow (\lambda v. x \overset{\gamma_2}{\leadsto_0^\star} v), \gamma''_2) \wedge$$
$$\forall x x'. f_V(x) = x' \Leftrightarrow x \overset{\gamma_2}{\longleftrightarrow} x'$$

Here, when we mention *mark* and *mark1*, the value 1s in the
original definition are changed to non-zero values.

Fig. 13. Proof sketch for bigraph copy - part 2

$$\frac{\displaystyle 100 \mapsto 42, 100, 0 \;\vdash\; 100 \mapsto 42, 100, 0 \uplus \mathrm{graph}_T(100, \hat{\gamma})}{\displaystyle \frac{100 \mapsto 42, 100, 0 \;\vdash\; \hat{\gamma}(100) = (42, 100, 0) \;\wedge\; 100 \mapsto 42, 100, 0 \uplus \mathrm{graph}_T(100, \hat{\gamma}) \uplus \mathrm{graph}_T(0, \hat{\gamma})}{100 \mapsto 42, 100, 0 \;\vdash\; \mathrm{graph}_T(100, \hat{\gamma})} (1)} (2)$$

(1) Unfold graph_T, dismiss first disjunct (contradiction), introduce existentials (which must be 42,100,0)

(2) simplify using $P * \mathrm{emp} \dashv\vdash P$ and remove pure conjunct

Fig. 14. An attempt to prove a "simple" entailment